

# Air Quality

November 08, 2021

## 1 IBM Supervised Learning: Regression peer- reviewed Project

by *Marwan Khalil*

### 1.1 Introduction

Air pollution will endanger human health and life in big cities, especially to the elderly and children. This is not an individual problem of one person but a global problem. Therefore, many countries in the world made air pollution monitoring and control stations in many cities to observe air pollutants such as NO<sub>2</sub>, CO, SO<sub>2</sub>, PM<sub>2.5</sub>, and PM<sub>10</sub> to alert the citizens about pollution index which exceeds the quality threshold.

Particulate Matter PM 2.5 is a fine atmospheric pollutant that has a diameter of fewer than 2.5 micrometers. Particulate Matter PM<sub>10</sub> is a coarse particulate that is 10 micrometers or less in diameter. Carbon Monoxide CO is a product of combustion of fuel such as coal, wood, or natural gas. Vehicular emission contributes to the majority of carbon monoxide let into our atmosphere. Nitrogen dioxide or nitrogen oxide expelled from high-temperature combustion: sulfur dioxide SO<sub>2</sub> and Sulphur Oxides SO produced by volcanoes and in industrial processes. Petroleum and Coal often contain sulfur compounds, and their combustion generates sulfur dioxide. Air pollution is caused by the presence of poison gases and substances; therefore, it is impacted by the meteorological factors of a particular place, such as temperature, humidity, rain, and wind.

In this project, the regression analysis technique is used to evaluate the relationship between these factors and **predict** nitrogen oxides NO<sub>x</sub> based on other parameters; **main objective of the analysis is to focus on prediction.**

### 1.2 Dataset description

The dataset contains 9358 instances of hourly averaged responses from an array of 5 metal oxide chemical sensors embedded in an Air Quality Chemical Multisensor Device. The device was located on the field in a significantly polluted area, at road level, within an Italian city. Data were recorded from March 2004 to February 2005 (one year) representing the longest freely available recordings of on field deployed air quality chemical sensor devices responses. Ground Truth hourly averaged concentrations for CO, Non Metanic Hydrocarbons, Benzene, Total Nitrogen Oxides (NO<sub>x</sub>) and Nitrogen Dioxide (NO<sub>2</sub>) and were provided by a co-located reference certified analyzer.

Dataset attribute information:

- Date (DD/MM/YYYY)
- Time (HH.MM.SS)
- True hourly averaged concentration CO in mg/m<sup>3</sup> (reference analyzer)

- PT08.S1 (tin oxide) hourly averaged sensor response (nominally CO targeted)
- True hourly averaged overall Non Metanic HydroCarbons concentration in microg/m<sup>3</sup> (reference analyzer)
- True hourly averaged Benzene concentration in microg/m<sup>3</sup> (reference analyzer)
- PT08.S2 (titania) hourly averaged sensor response (nominally NMHC targeted)
- True hourly averaged NOx concentration in ppb (reference analyzer)
- PT08.S3 (tungsten oxide) hourly averaged sensor response (nominally NOx targeted)
- True hourly averaged NO2 concentration in microg/m<sup>3</sup> (reference analyzer)
- PT08.S4 (tungsten oxide) hourly averaged sensor response (nominally NO2 targeted)
- PT08.S5 (indium oxide) hourly averaged sensor response (nominally O3 targeted)
- Temperature in Â°C
- Relative Humidity (%)
- AH Absolute Humidity

This dataset is from the UCI machine learning repository and contains hourly averaged responses from an air quality multi-sensor device that was located in a significantly polluted area at road level in an undisclosed Italian city. This data was collected over the course of approx one year (from March 2004 - February 2005)

For reference: <https://archive.ics.uci.edu/ml/datasets/Air+Quality#>

*Note:* Evidences of cross-sensitivities as well as both concept and sensor drifts are present as described in De Vito et al., Sens. And Act. B, Vol. 129,2,2008 (citation required) eventually affecting sensors concentration estimation capabilities. **Missing values are tagged with -200 value.**

```
[1]: # importing dependencies

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings, pprint

from sklearn import metrics
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.model_selection import KFold, cross_val_predict, train_test_split
from sklearn.linear_model import LinearRegression, Lasso, Ridge, RidgeCV,
↳LassoCV, ElasticNetCV
from sklearn.pipeline import Pipeline

%matplotlib inline

print("All relevant modules were imported.")
```

All relevant modules were imported.

### 1.3 Data cleaning and feature engineering.

```
[2]: airq = pd.read_csv('AirQualityUCI.csv', sep=',', delimiter=";", decimal=",")
```

```
[3]: airq.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9471 entries, 0 to 9470
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Date                  9357 non-null  object 
1   Time                  9357 non-null  object 
2   CO(GT)                9357 non-null  float64
3   PT08.S1(CO)          9357 non-null  float64
4   NMHC(GT)              9357 non-null  float64
5   C6H6(GT)              9357 non-null  float64
6   PT08.S2(NMHC)         9357 non-null  float64
7   NOx(GT)               9357 non-null  float64
8   PT08.S3(NOx)          9357 non-null  float64
9   NO2(GT)               9357 non-null  float64
10  PT08.S4(NO2)          9357 non-null  float64
11  PT08.S5(O3)           9357 non-null  float64
12  T                     9357 non-null  float64
13  RH                    9357 non-null  float64
14  AH                    9357 non-null  float64
15  Unnamed: 15           0 non-null     float64
16  Unnamed: 16           0 non-null     float64
dtypes: float64(15), object(2)
memory usage: 1.2+ MB
```

```
[4]: airq.columns
```

```
[4]: Index(['Date', 'Time', 'CO(GT)', 'PT08.S1(CO)', 'NMHC(GT)', 'C6H6(GT)',
          'PT08.S2(NMHC)', 'NOx(GT)', 'PT08.S3(NOx)', 'NO2(GT)', 'PT08.S4(NO2)',
          'PT08.S5(O3)', 'T', 'RH', 'AH', 'Unnamed: 15', 'Unnamed: 16'],
          dtype='object')
```

```
[5]: airq.describe()
```

```
[5]:
```

	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	\
count	9357.000000	9357.000000	9357.000000	9357.000000	9357.000000	
mean	-34.207524	1048.990061	-159.090093	1.865683	894.595276	
std	77.657170	329.832710	139.789093	41.380206	342.333252	
min	-200.000000	-200.000000	-200.000000	-200.000000	-200.000000	
25%	0.600000	921.000000	-200.000000	4.000000	711.000000	
50%	1.500000	1053.000000	-200.000000	7.900000	895.000000	
75%	2.600000	1221.000000	-200.000000	13.600000	1105.000000	

max	11.900000	2040.000000	1189.000000	63.700000	2214.000000
-----	-----------	-------------	-------------	-----------	-------------

	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3) \
count	9357.000000	9357.000000	9357.000000	9357.000000	9357.000000
mean	168.616971	794.990168	58.148873	1391.479641	975.072032
std	257.433866	321.993552	126.940455	467.210125	456.938184
min	-200.000000	-200.000000	-200.000000	-200.000000	-200.000000
25%	50.000000	637.000000	53.000000	1185.000000	700.000000
50%	141.000000	794.000000	96.000000	1446.000000	942.000000
75%	284.000000	960.000000	133.000000	1662.000000	1255.000000
max	1479.000000	2683.000000	340.000000	2775.000000	2523.000000


  

	T	RH	AH	Unnamed: 15	Unnamed: 16
count	9357.000000	9357.000000	9357.000000	0.0	0.0
mean	9.778305	39.485380	-6.837604	NaN	NaN
std	43.203623	51.216145	38.976670	NaN	NaN
min	-200.000000	-200.000000	-200.000000	NaN	NaN
25%	10.900000	34.100000	0.692300	NaN	NaN
50%	17.200000	48.600000	0.976800	NaN	NaN
75%	24.100000	61.900000	1.296200	NaN	NaN
max	44.600000	88.700000	2.231000	NaN	NaN

```
[6]: plt.figure(figsize=(9,7))
plt.title('Missing values', fontweight='bold')

ax = sns.heatmap(airq.isnull(),yticklabels=False,cbar='viridis',cmap='viridis')

plt.show()
```



```
[7]: airq.isnull().sum()
```

```
[7]: Date          114
     Time          114
     CO(GT)        114
     PT08.S1(CO)    114
     NMHC(GT)       114
     C6H6(GT)       114
     PT08.S2(NMHC)  114
     NOx(GT)        114
     PT08.S3(NOx)   114
     NO2(GT)        114
```

```

PT08.S4(NO2)      114
PT08.S5(O3)       114
T                 114
RH                114
AH                114
Unnamed: 15       9471
Unnamed: 16       9471
dtype: int64

```

```
[8]: airq = airq.drop(["Unnamed: 15", "Unnamed: 16"], axis=1)
```

```
[9]: airq.dropna(inplace=True)
```

```
[10]: airq.set_index("Date", inplace=True)
```

```
[11]: airq.index = pd.to_datetime(airq.index)
```

```
[12]: type(airq.index)
```

```
[12]: pandas.core.indexes.datetimes.DatetimeIndex
```

```
[13]: airq['Time'] = pd.to_datetime(airq['Time'], format= '%H.%M.%S').dt.hour
```

```
[14]: airq.apply(lambda x : x == -200).sum()
```

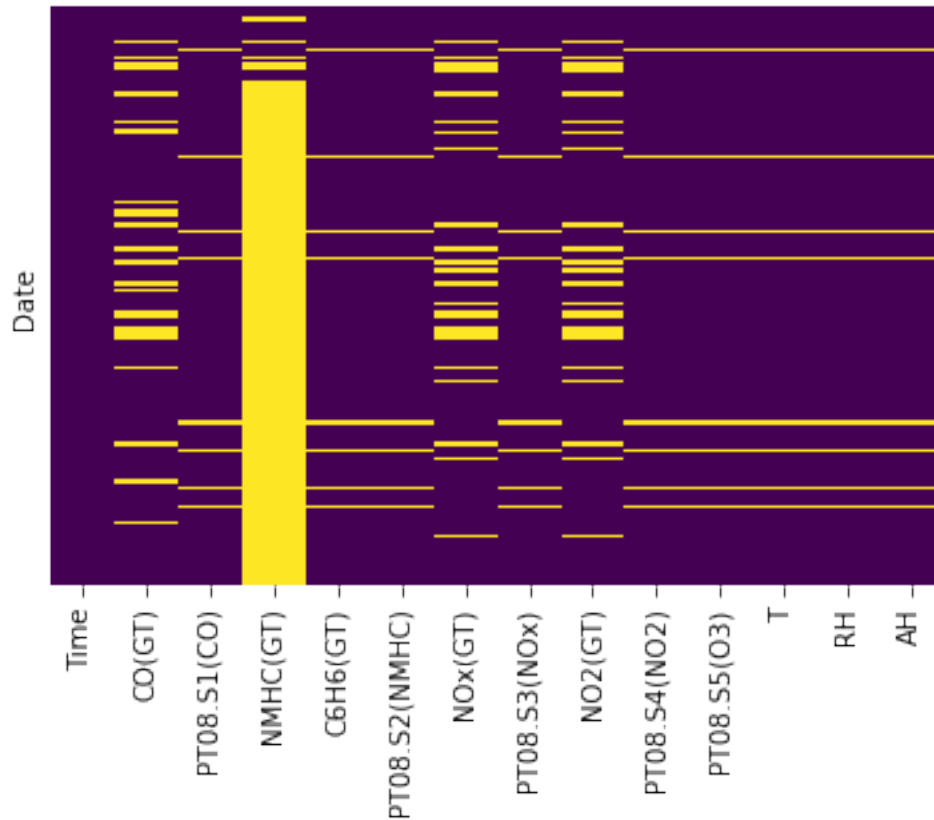
```

[14]: Time                0
      CO(GT)             1683
      PT08.S1(CO)         366
      NMHC(GT)           8443
      C6H6(GT)           366
      PT08.S2(NMHC)       366
      NOx(GT)            1639
      PT08.S3(NOx)       366
      NO2(GT)            1642
      PT08.S4(NO2)       366
      PT08.S5(O3)       366
      T                 366
      RH                366
      AH                366
      dtype: int64

```

```
[15]: sns.heatmap(airq.isin([-200]), yticklabels=False, cbar=False, cmap='viridis')
```

```
[15]: <AxesSubplot:ylabel='Date'>
```

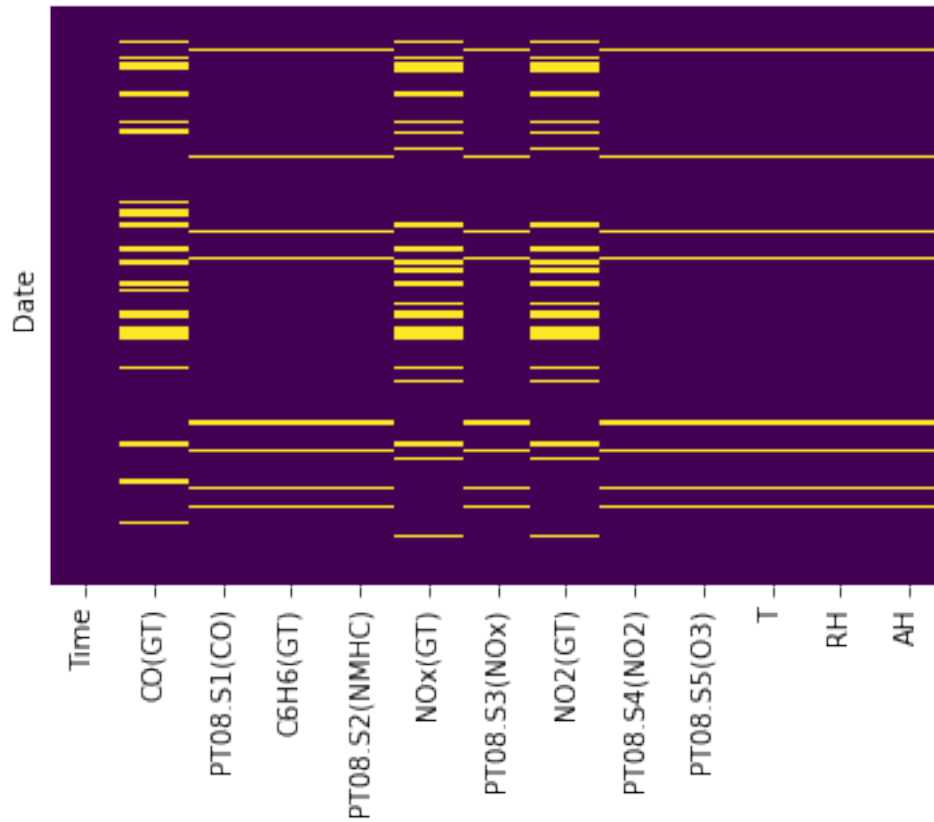


The NMHC(GT) column is missing many values, more than 85% values are NaN; 8443 out of 9357. Thus this column can be removed from the dataset as these values are very less likely to be of any significant importance in this dataset.

```
[16]: airq.drop('NMHC(GT)', axis=1, inplace=True)
```

```
[17]: sns.heatmap(airq.isin([-200]), yticklabels=False, cbar=False, cmap='viridis')
```

```
[17]: <AxesSubplot:ylabel='Date'>
```



```
[18]: airq.isin([-200]).sum()
```

```
[18]: Time          0
      CO(GT)       1683
      PT08.S1(CO)   366
      C6H6(GT)      366
      PT08.S2(NMHC) 366
      NOx(GT)       1639
      PT08.S3(NOx)   366
      NO2(GT)       1642
      PT08.S4(NO2)   366
      PT08.S5(O3)    366
      T             366
      RH            366
      AH            366
      dtype: int64
```

```
[19]: airq.replace(to_replace= -200, value= np.NaN, inplace= True)
```



```
[20]: # Fill NaN values with average of particular date
```

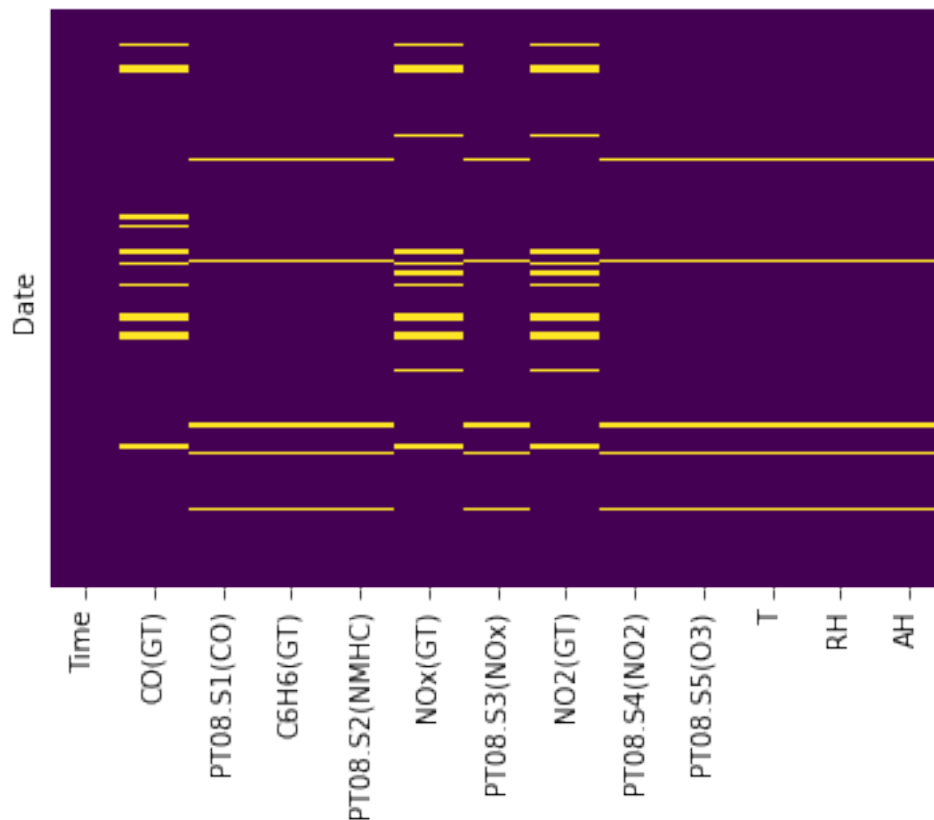
```
def remove_outlier(col):  
    airq[col] = airq.groupby('Date')[col].transform(lambda x: x.fillna(x.  
    ↪mean()))
```

```
[21]: col_list = airq.columns[1:]
```

```
for i in col_list:  
    remove_outlier(i)
```

```
[22]: sns.heatmap(airq.isna(),yticklabels=False,cbar=False,cmap='viridis')
```

```
[22]: <AxesSubplot:ylabel='Date'>
```




```
[23]: # forward fill method for removing the leftover nan values
```

```
airq.fillna(method='ffill', inplace= True)
```

```
[24]: sns.heatmap(airq.isna(),yticklabels=False,cbar=False,cmap='viridis')
```

```
[24]: <AxesSubplot:ylabel='Date'>
```



```
[25]: airq.isnull().any()
```

```
[25]: Time                False
      CO(GT)              False
      PT08.S1(CO)         False
      C6H6(GT)            False
      PT08.S2(NMHC)       False
      NOx(GT)             False
      PT08.S3(NOx)        False
      NO2(GT)             False
      PT08.S4(NO2)        False
      PT08.S5(O3)         False
      T                   False
      RH                  False
      AH                  False
      dtype: bool
```

## 1.4 EXPLORATORY DATA ANALYSIS

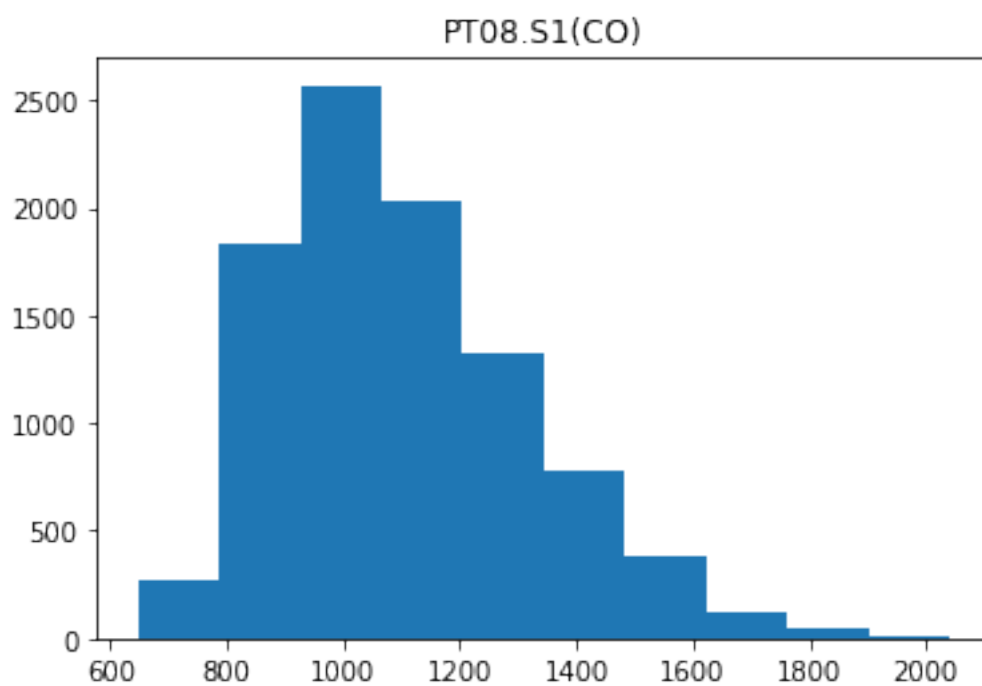
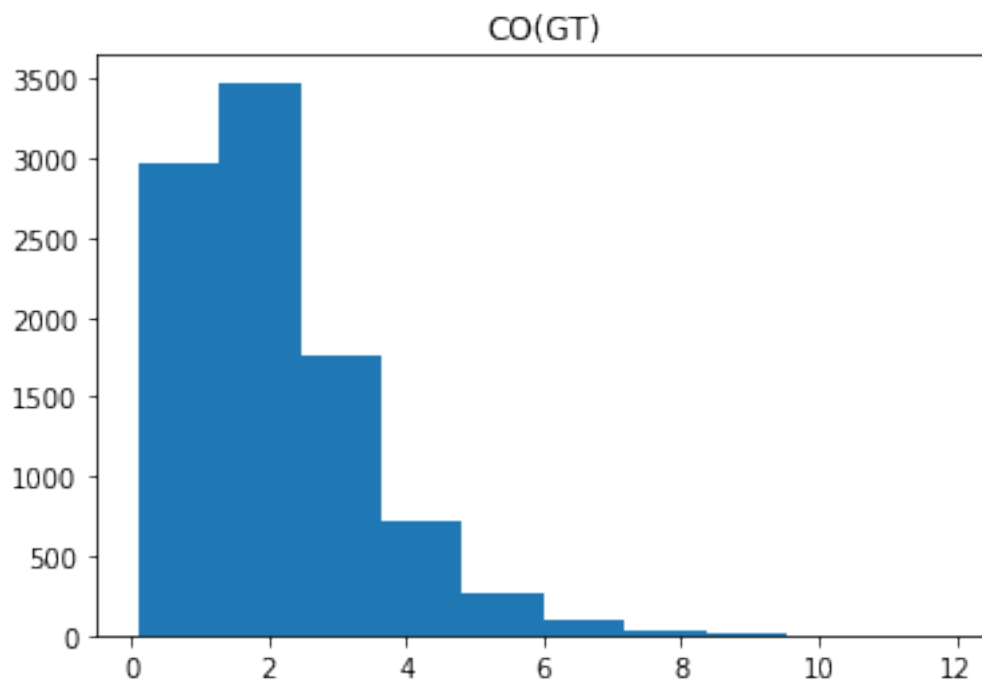
```
[26]: airq.dtypes.value_counts()
```

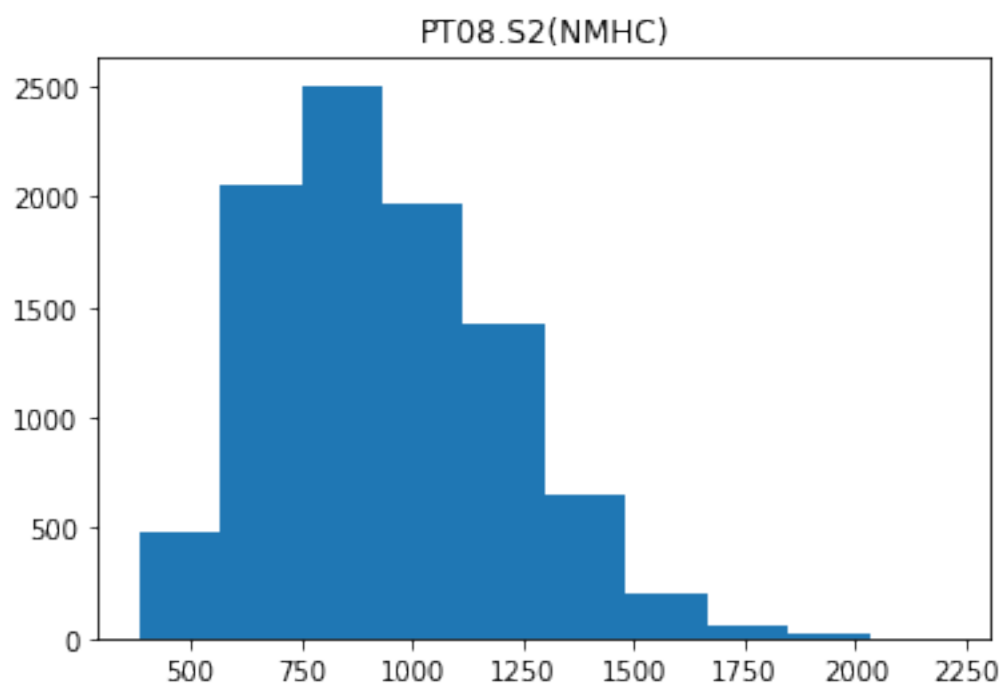
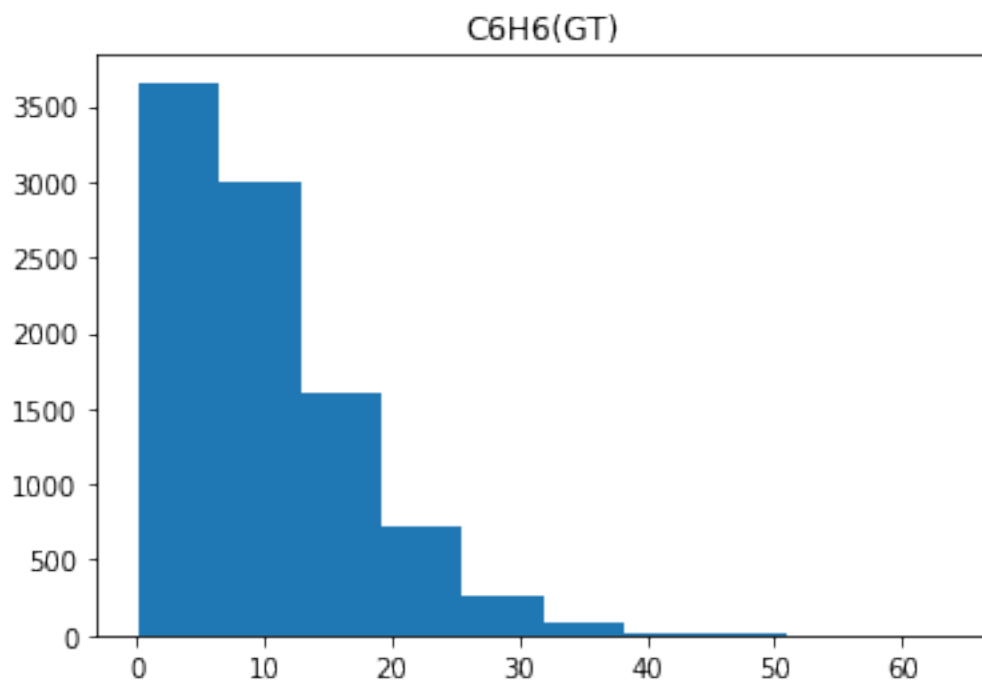
```
[26]: float64    12  
      int64      1  
      dtype: int64
```

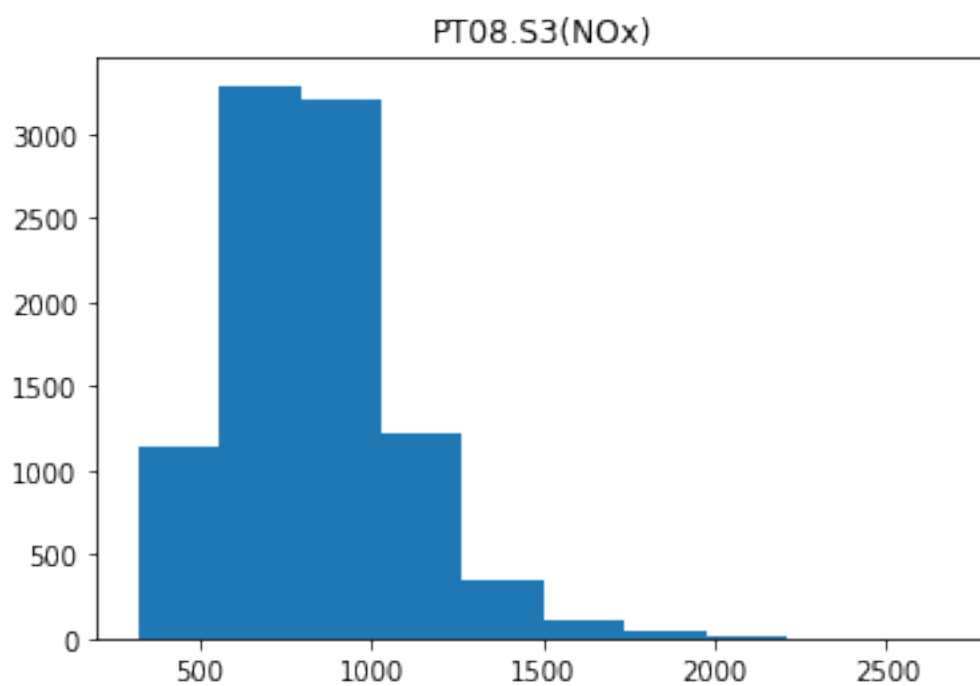
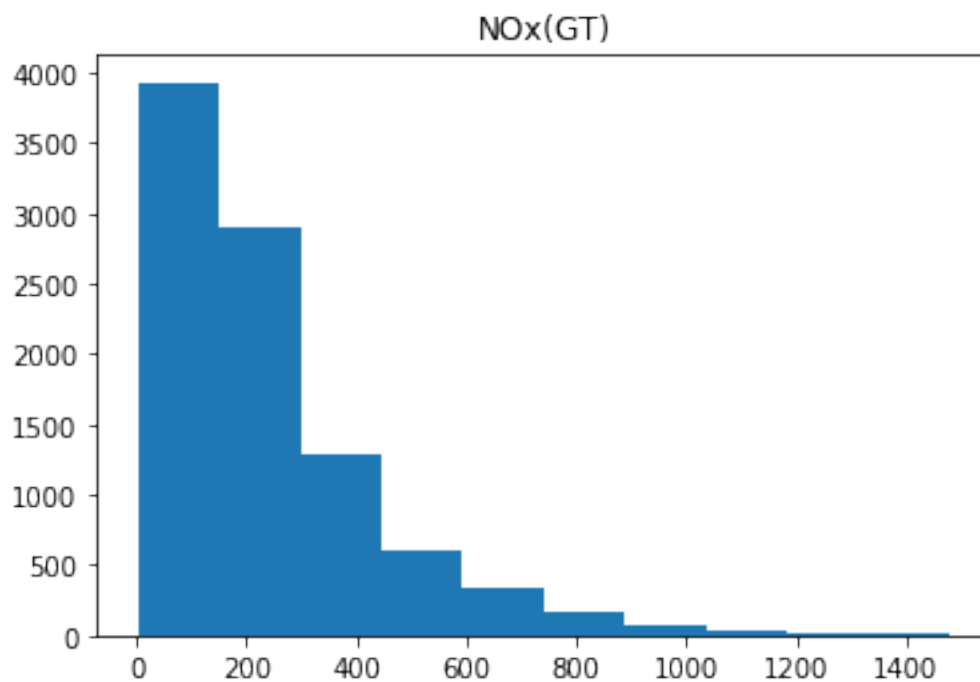
```
[27]: columns = list(airq.columns)  
  
      columns[1:]
```

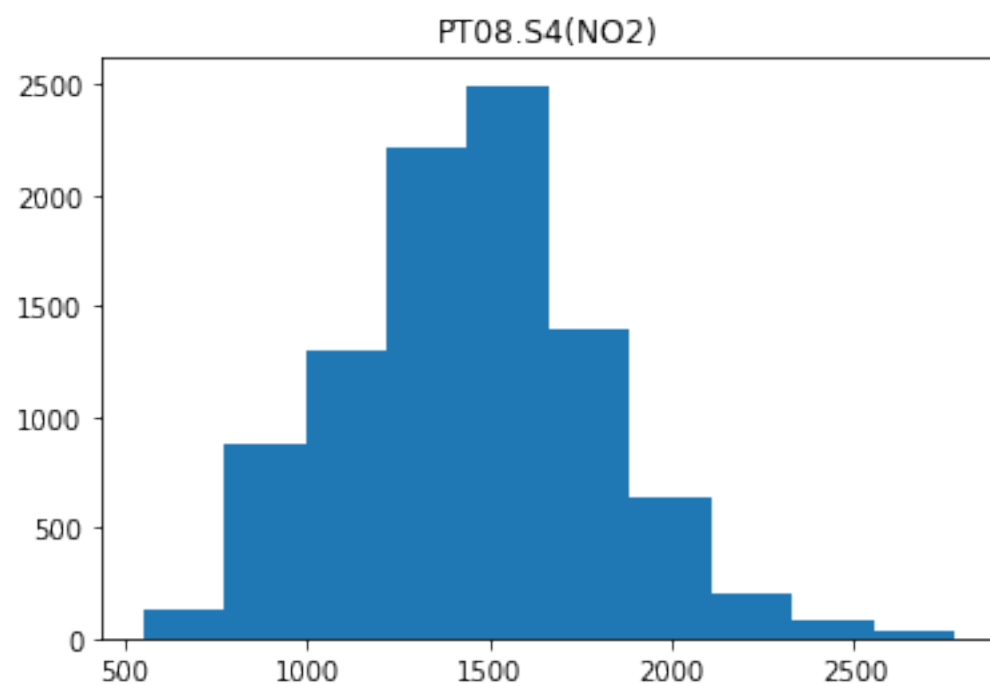
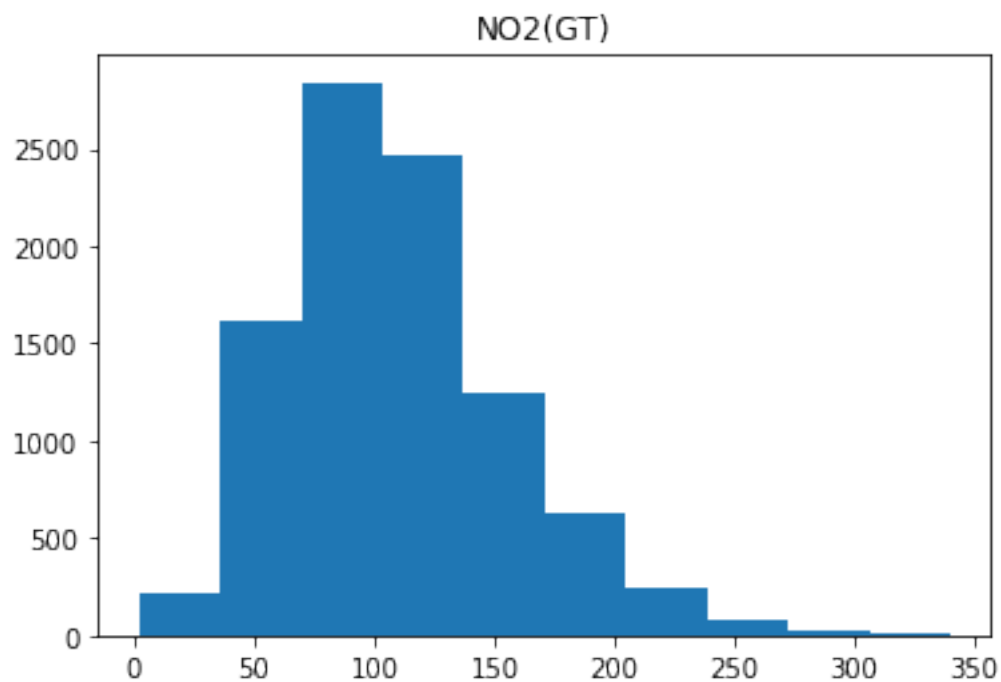
```
[27]: ['CO(GT) ',  
      'PT08.S1(CO) ',  
      'C6H6(GT) ',  
      'PT08.S2(NMHC) ',  
      'NOx(GT) ',  
      'PT08.S3(NOx) ',  
      'NO2(GT) ',  
      'PT08.S4(NO2) ',  
      'PT08.S5(O3) ',  
      'T',  
      'RH',  
      'AH']
```

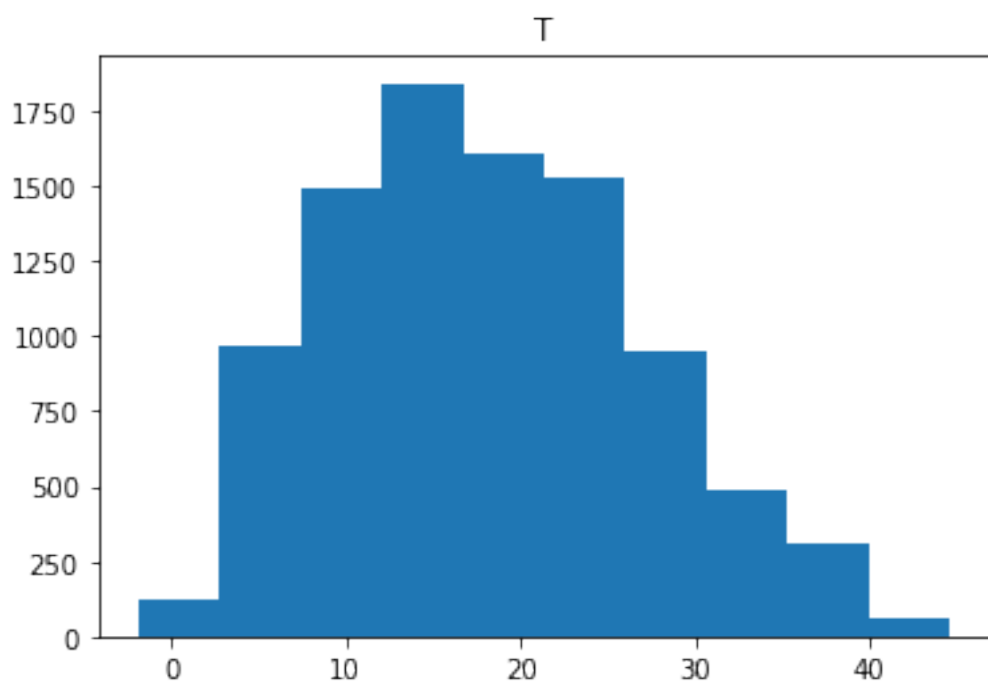
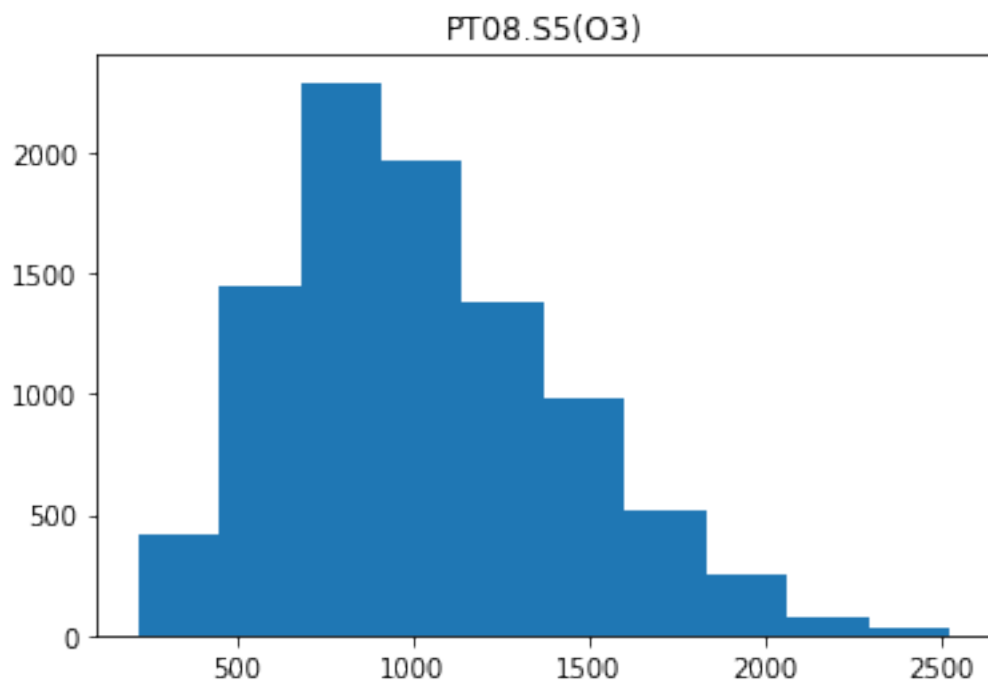
```
[28]: for col in columns[1:]:  
      plt.hist(airq[col])  
      plt.title(col)  
      plt.show()
```




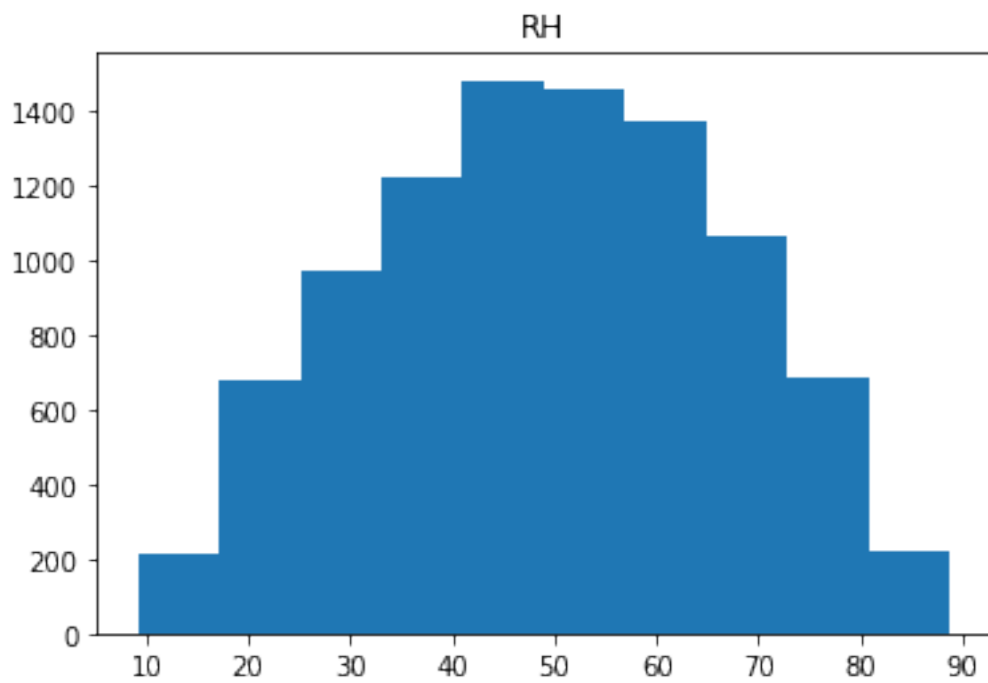








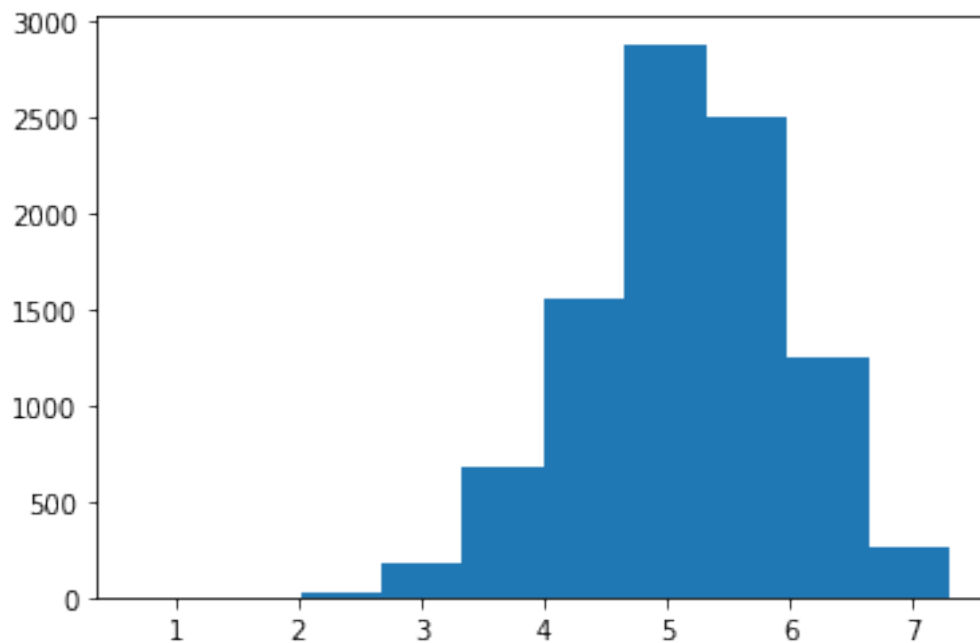




```
[29]: # Log-transform the skewed features
```


```
plt.hist(airq['NOx(GT)'].apply(lambda x: np.log(x)))
```

```
[29]: (array([1.000e+00, 3.000e+00, 2.300e+01, 1.850e+02, 6.860e+02, 1.556e+03,  
          2.882e+03, 2.504e+03, 1.255e+03, 2.620e+02]),  
      array([0.69314718, 1.35374461, 2.01434204, 2.67493947, 3.33553689,  
          3.99613432, 4.65673175, 5.31732918, 5.97792661, 6.63852403,  
          7.29912146]),  
      <BarContainer object of 10 artists>)
```

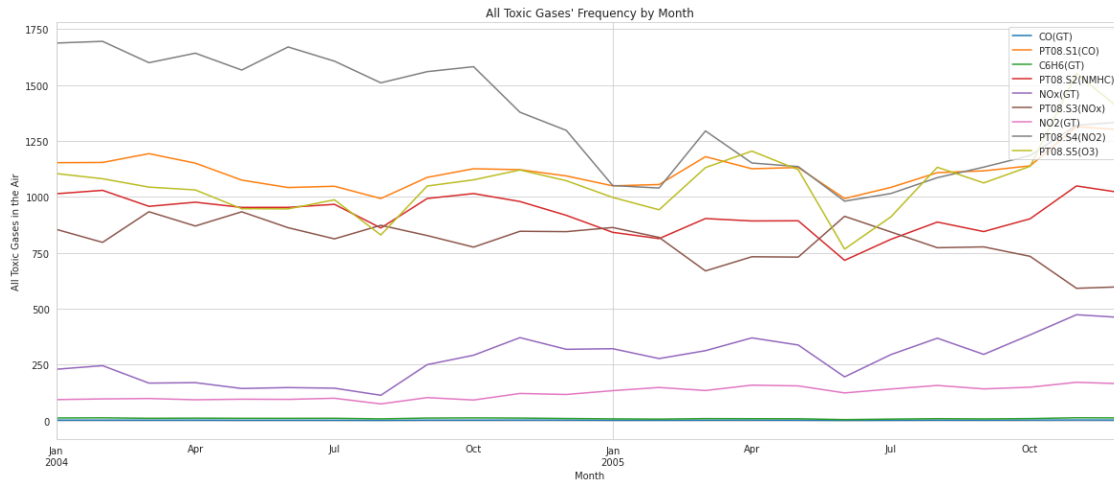


```
[30]: sns.set_style('whitegrid')  
eda_air = airq.drop(['Time', 'RH', 'AH', 'T'], axis=1)  
sns.pairplot(eda_air)
```

```
[30]: <seaborn.axisgrid.PairGrid at 0x7fbcce527b20>
```

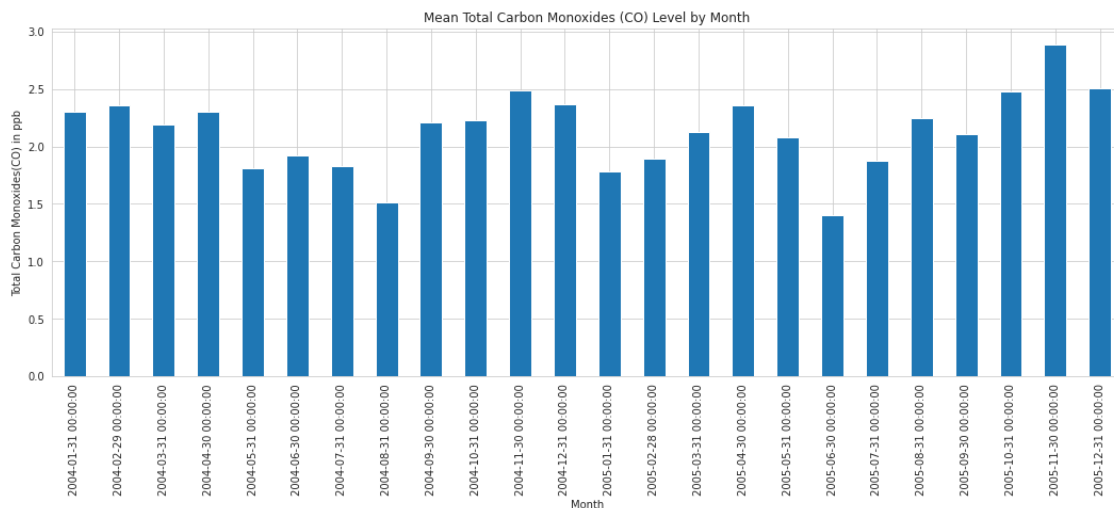


```
[31]: airq.drop(['Time', 'RH', 'AH', 'T'], axis=1).resample('M').mean().plot(figsize = (20,8))
plt.legend(loc=1)
plt.xlabel('Month')
plt.ylabel('All Toxic Gases in the Air')
plt.title("All Toxic Gases' Frequency by Month");
```



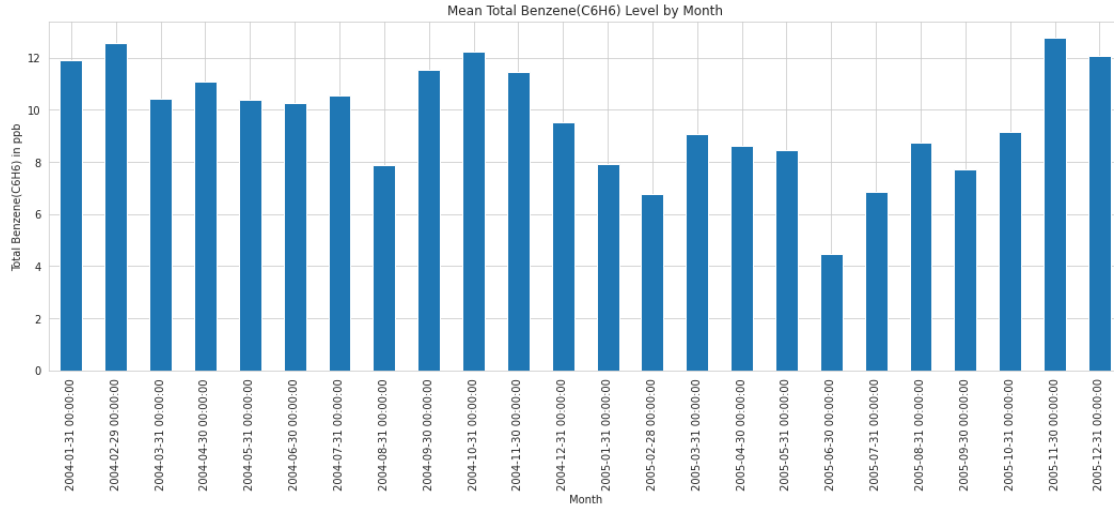
```
[32]: airq['CO(GT)'].resample('M').mean().plot(kind='bar', figsize=(18,6))
plt.xlabel('Month')
plt.ylabel('Total Carbon Monoxides(CO) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Carbon Monoxides (CO) Level by Month")
```

[32]: Text(0.5, 1.0, 'Mean Total Carbon Monoxides (CO) Level by Month')



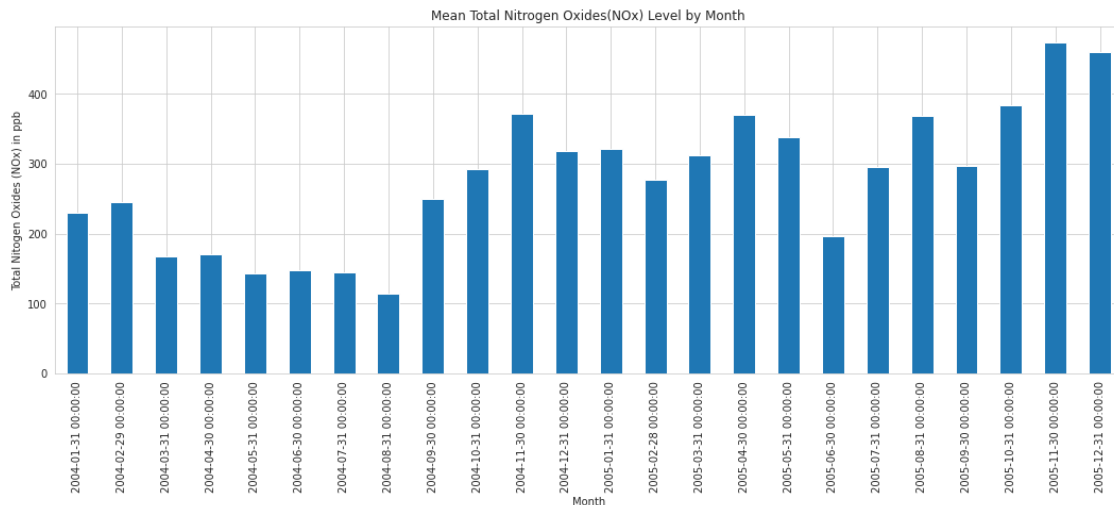
```
[33]: airq['C6H6(GT)'].resample('M').mean().plot(kind='bar', figsize=(18,6))
plt.xlabel('Month')
plt.ylabel('Total Benzene(C6H6) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Benzene(C6H6) Level by Month")
```

[33]: Text(0.5, 1.0, 'Mean Total Benzene(C6H6) Level by Month')



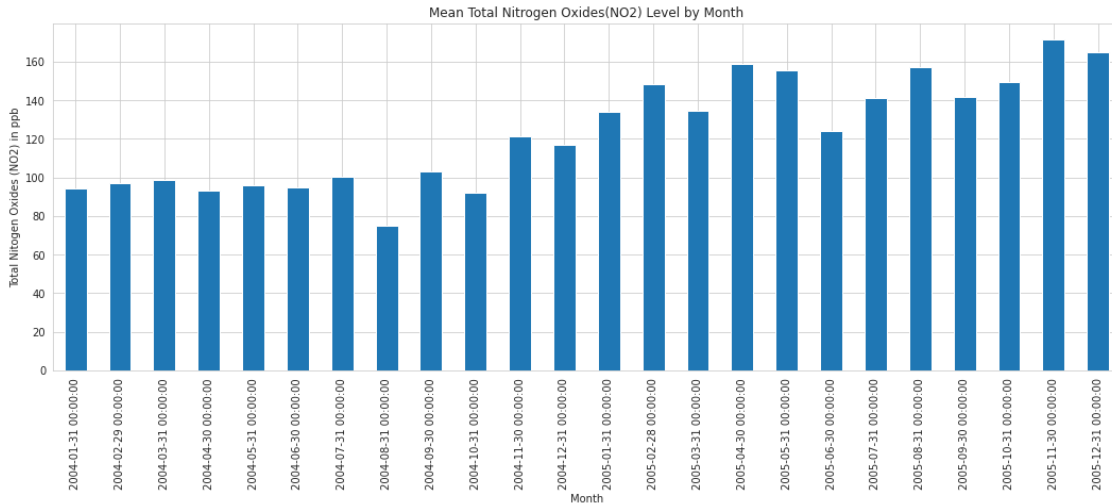
```
[34]: airq['NOx(GT)'].resample('M').mean().plot(kind='bar', figsize=(18,6))
plt.xlabel('Month')
plt.ylabel('Total Nitrogen Oxides (NOx) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Nitrogen Oxides(NOx) Level by Month")
```

[34]: Text(0.5, 1.0, 'Mean Total Nitrogen Oxides(NOx) Level by Month')



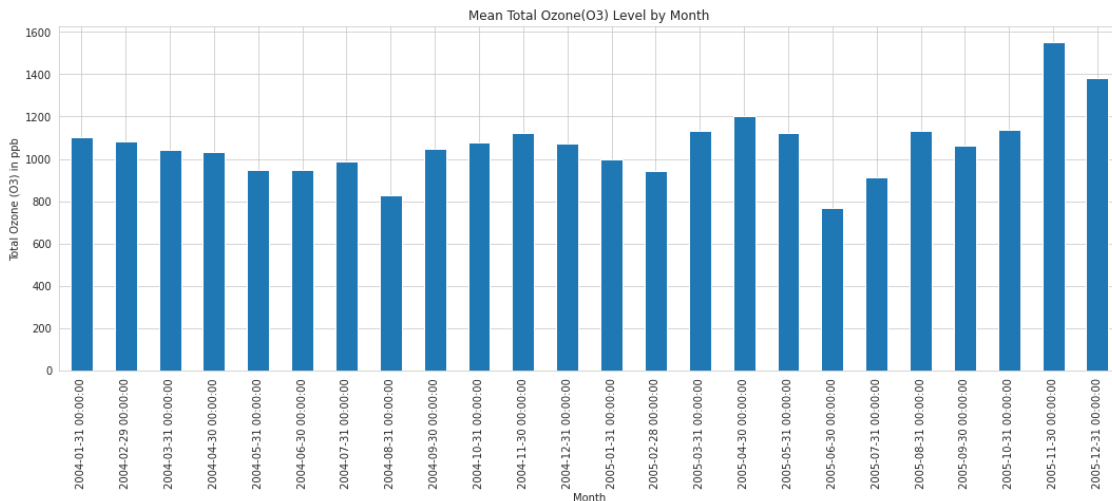
```
[35]: airq['NO2(GT)'].resample('M').mean().plot(kind='bar', figsize=(18,6))
plt.xlabel('Month')
plt.ylabel('Total Nitrogen Oxides (NO2) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Nitrogen Oxides(NO2) Level by Month")
```

[35]: Text(0.5, 1.0, 'Mean Total Nitrogen Oxides(NO2) Level by Month')



```
[36]: airq['PT08.S5(O3)'].resample('M').mean().plot(kind='bar', figsize=(18,6))
plt.xlabel('Month')
plt.ylabel('Total Ozone (O3) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Ozone(O3) Level by Month")
```

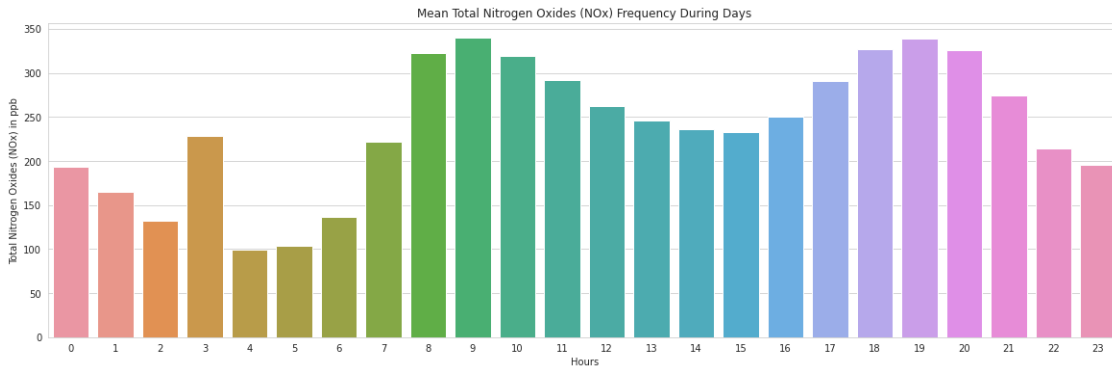
[36]: Text(0.5, 1.0, 'Mean Total Ozone(O3) Level by Month')



```
[37]: plt.figure(figsize=(20,6))
sns.barplot(x='Time',y='NOx(GT)',data=airq, ci=False)
plt.xlabel('Hours')
```

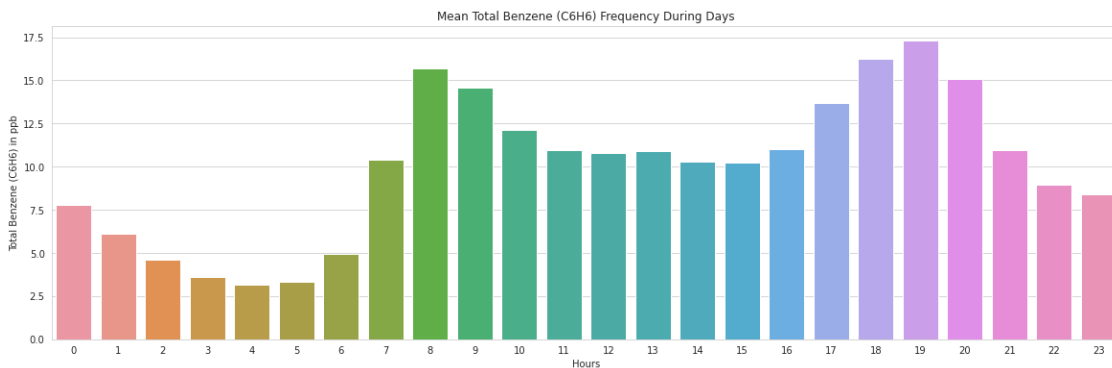
```
plt.ylabel('Total Nitrogen Oxides (NOx) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Nitrogen Oxides (NOx) Frequency During Days")
```

[37]: Text(0.5, 1.0, 'Mean Total Nitrogen Oxides (NOx) Frequency During Days')




```
[38]: plt.figure(figsize=(20,6))
sns.barplot(x='Time',y='C6H6(GT)',data=airq, ci=False)
plt.xlabel('Hours')
plt.ylabel('Total Benzene (C6H6) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Benzene (C6H6) Frequency During Days")
```

[38]: Text(0.5, 1.0, 'Mean Total Benzene (C6H6) Frequency During Days')



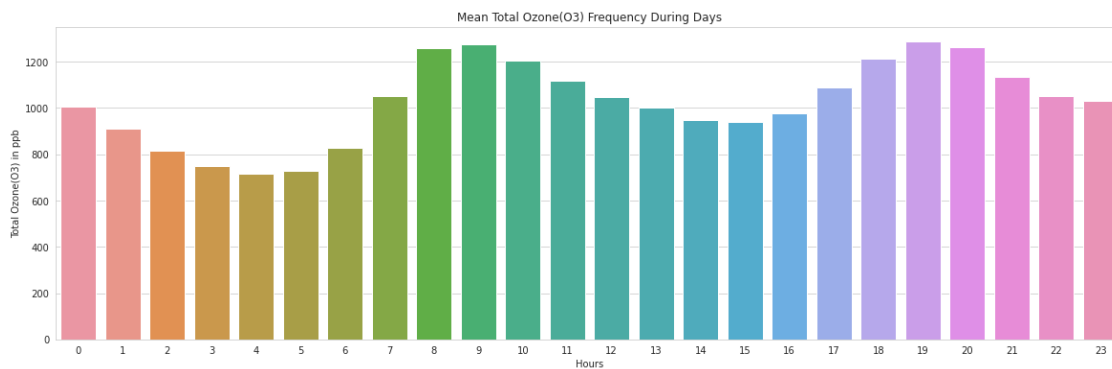
```
[39]: plt.figure(figsize=(20,6))
sns.barplot(x='Time',y='NO2(GT)',data=airq, ci=False)
plt.xlabel('Hours')
plt.ylabel('Total Nitrogen Oxides(NO2) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Nitrogen Oxides(NO2) Frequency During Days")
```

[39]: Text(0.5, 1.0, 'Mean Total Nitrogen Oxides(NO2) Frequency During Days')



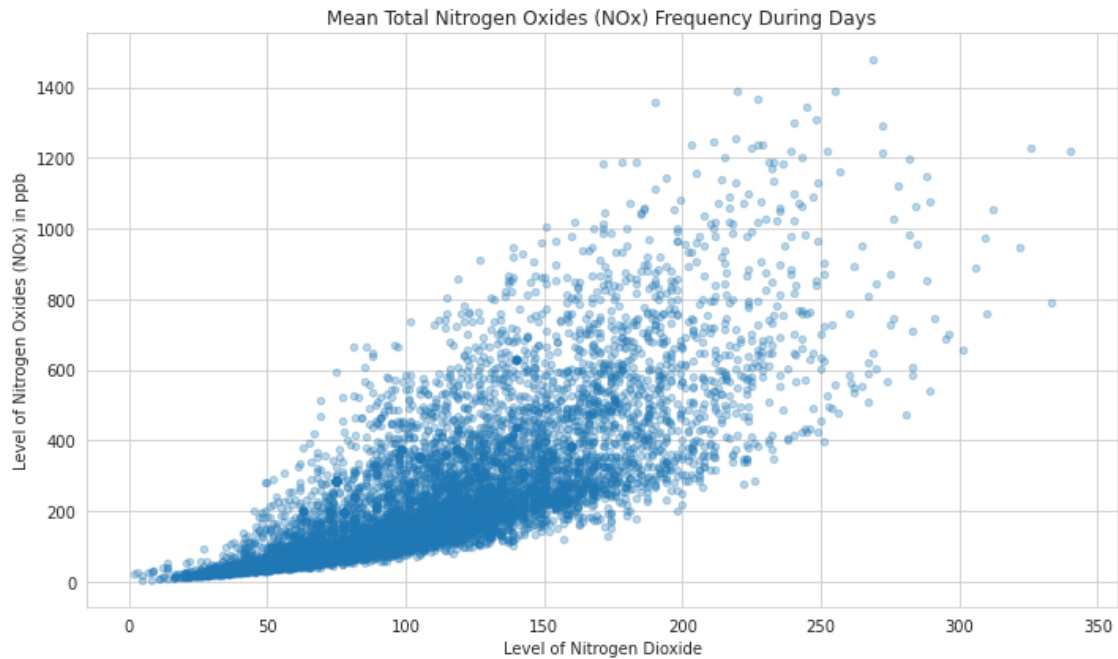
```
[40]: plt.figure(figsize=(20,6))
sns.barplot(x='Time',y='PT08.S5(O3)',data=airq, ci=False)
plt.xlabel('Hours')
plt.ylabel('Total Ozone(O3) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Ozone(O3) Frequency During Days")
```

[40]: Text(0.5, 1.0, 'Mean Total Ozone(O3) Frequency During Days')



```
[41]: airq.plot(x='NO2(GT)',y='NOx(GT)', kind='scatter', figsize = (10,6), alpha=0.3)
plt.xlabel('Level of Nitrogen Dioxide')
plt.ylabel('Level of Nitrogen Oxides (NOx) in ppb') # Parts per billion (ppb)
plt.title("Mean Total Nitrogen Oxides (NOx) Frequency During Days")
plt.tight_layout();
```






```
[42]: corrPearson = airq.corr(method="pearson")
```

```
plt.figure(figsize=(10,8))
sns.heatmap(corrPearson,
            annot=True,
            fmt=".2f",
            linewidth=.20,
            cmap='BrBG',
            vmin=-1,
            vmax=+1
            )
```

```
[42]: <AxesSubplot:>
```



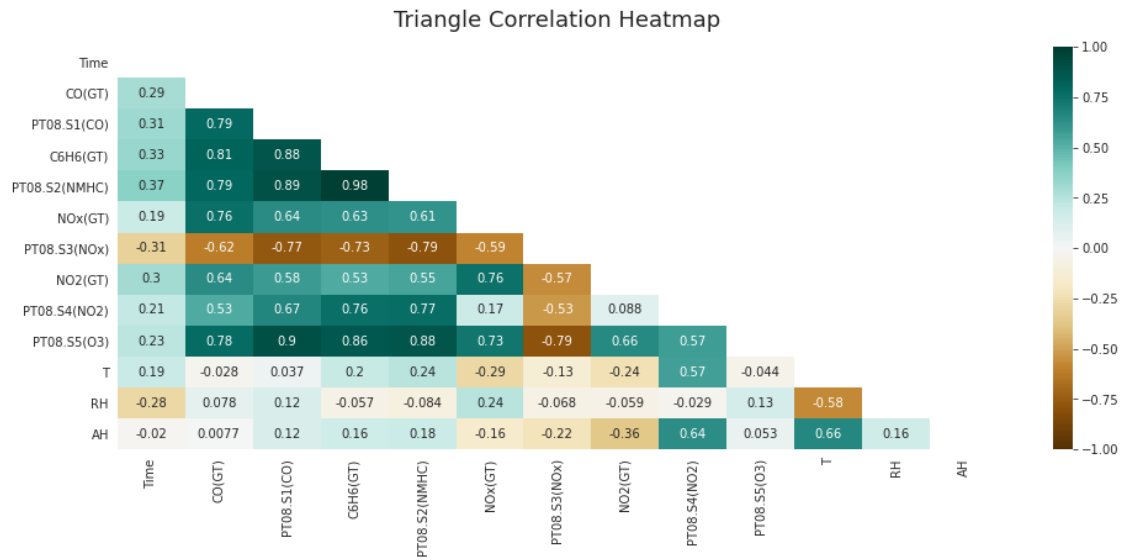
```
[43]: plt.figure(figsize=(16,6))

mask = np.triu(np.ones_like(corrPearson, dtype=np.bool))

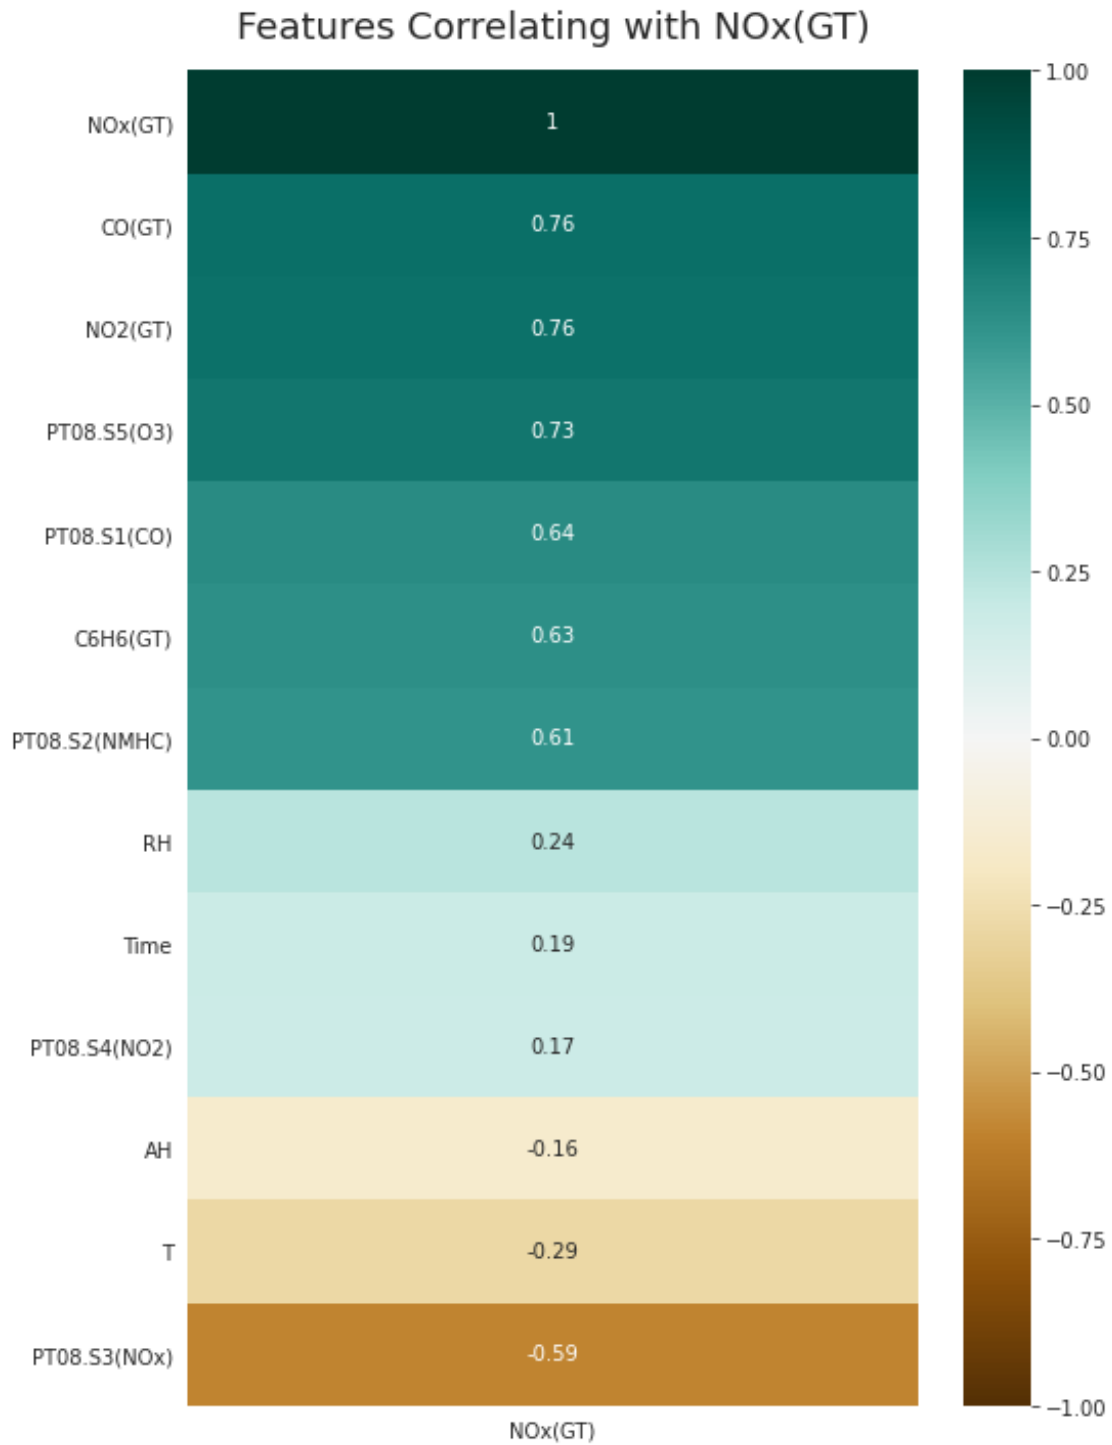
heatmap = sns.heatmap(corrPearson, mask=mask, vmin=-1, vmax=1, annot=True,
    cmap='BrBG')

heatmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18},
    pad=16);
```

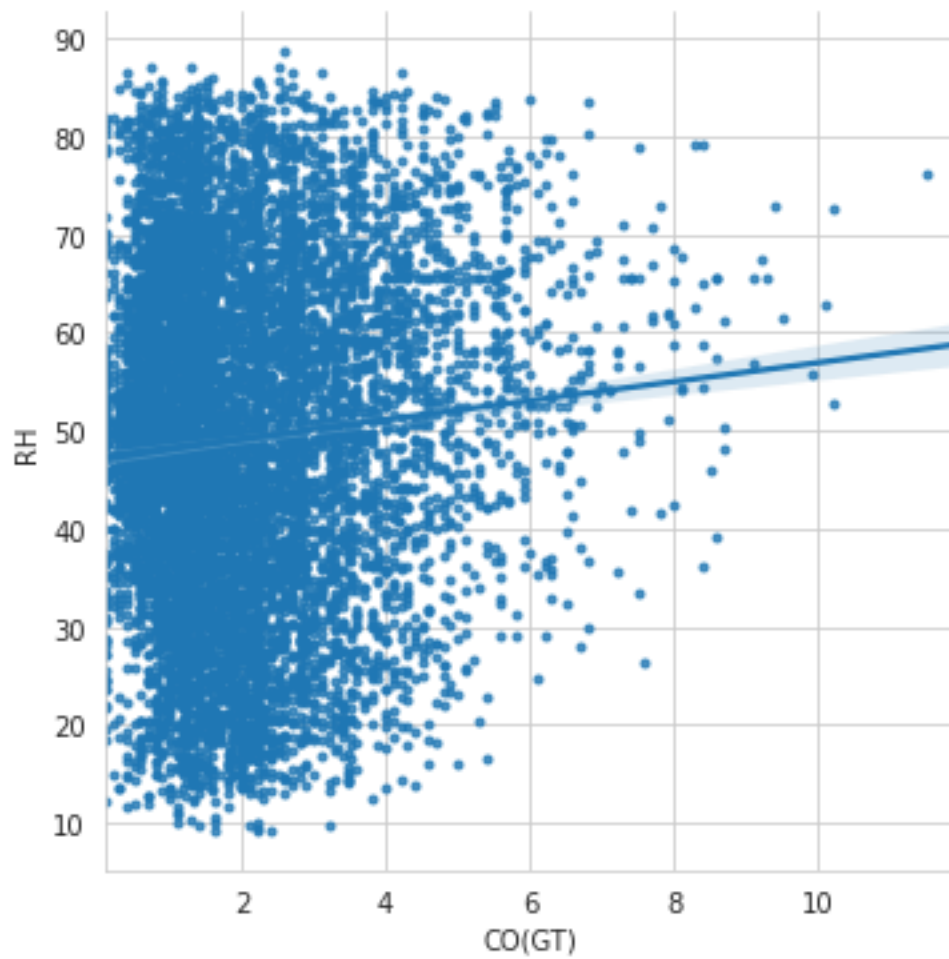
<ipython-input-43-b13df2c2b04c>:3: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool\_` here.  
 Deprecated in NumPy 1.20; for more details and guidance:  
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>  
 mask = np.triu(np.ones\_like(corrPearson, dtype=np.bool))





```
[44]: plt.figure(figsize=(8, 12))
heatmap = sns.heatmap(corrPearson[['NOx(GT)']].sort_values(by='NOx(GT)',
↪ascending=False), vmin=-1, vmax=1, annot=True, cmap='BrBG')
heatmap.set_title('Features Correlating with NOx(GT)', fontdict={'fontsize':
↪18}, pad=16);
```

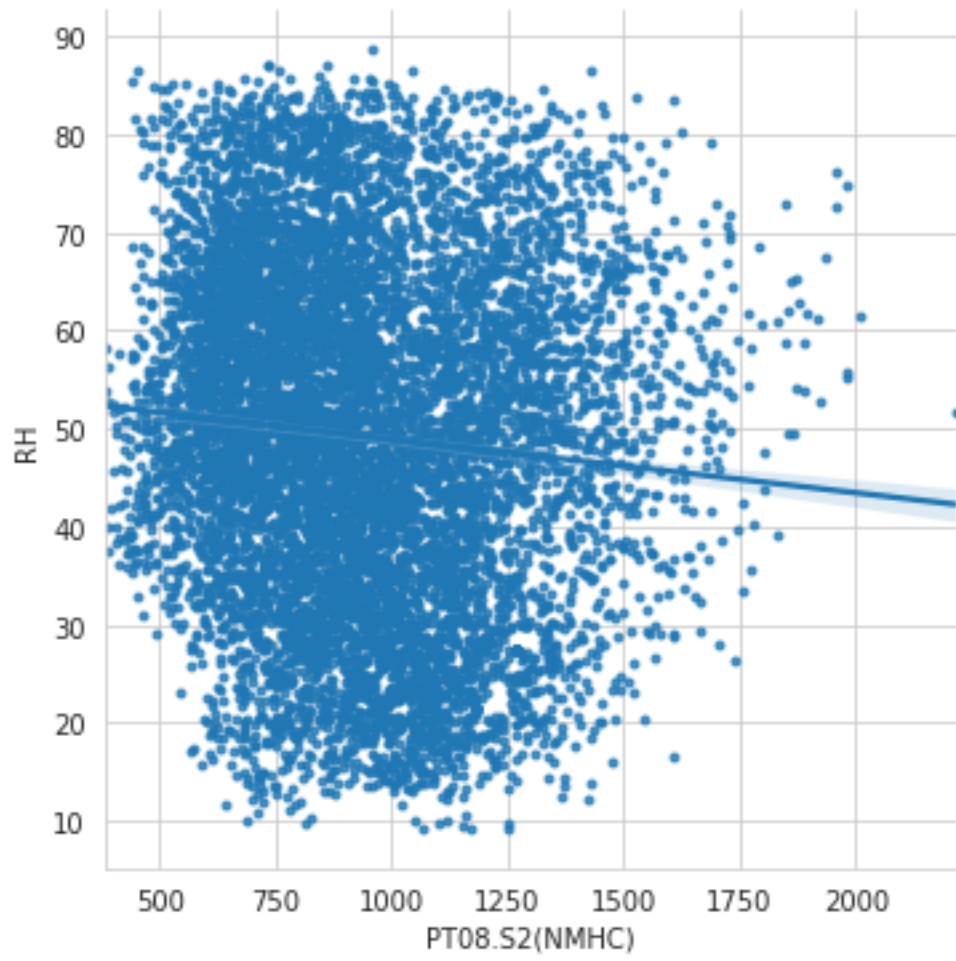


```
[45]: column_=airq.columns.tolist()[1:]
      for i in airq.columns.tolist()[1:]:
          sns.lmplot(x=i,y='RH',data=airq,markers='.')
```

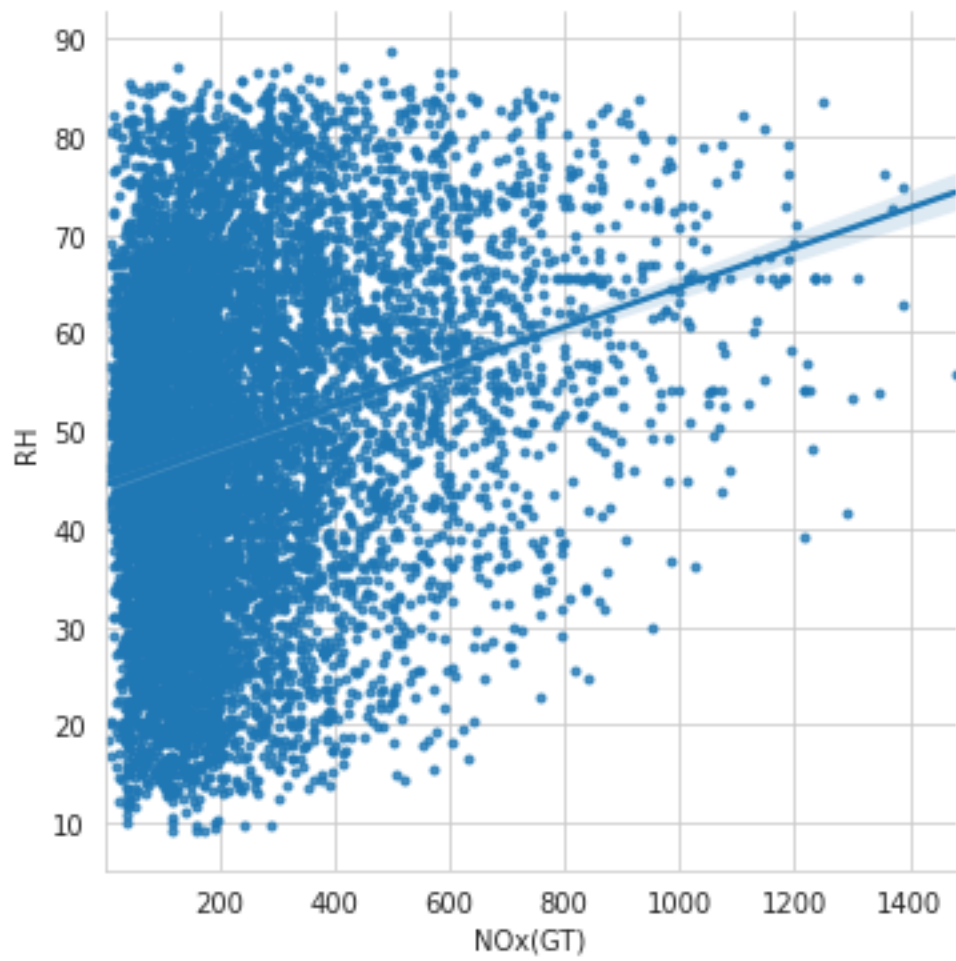


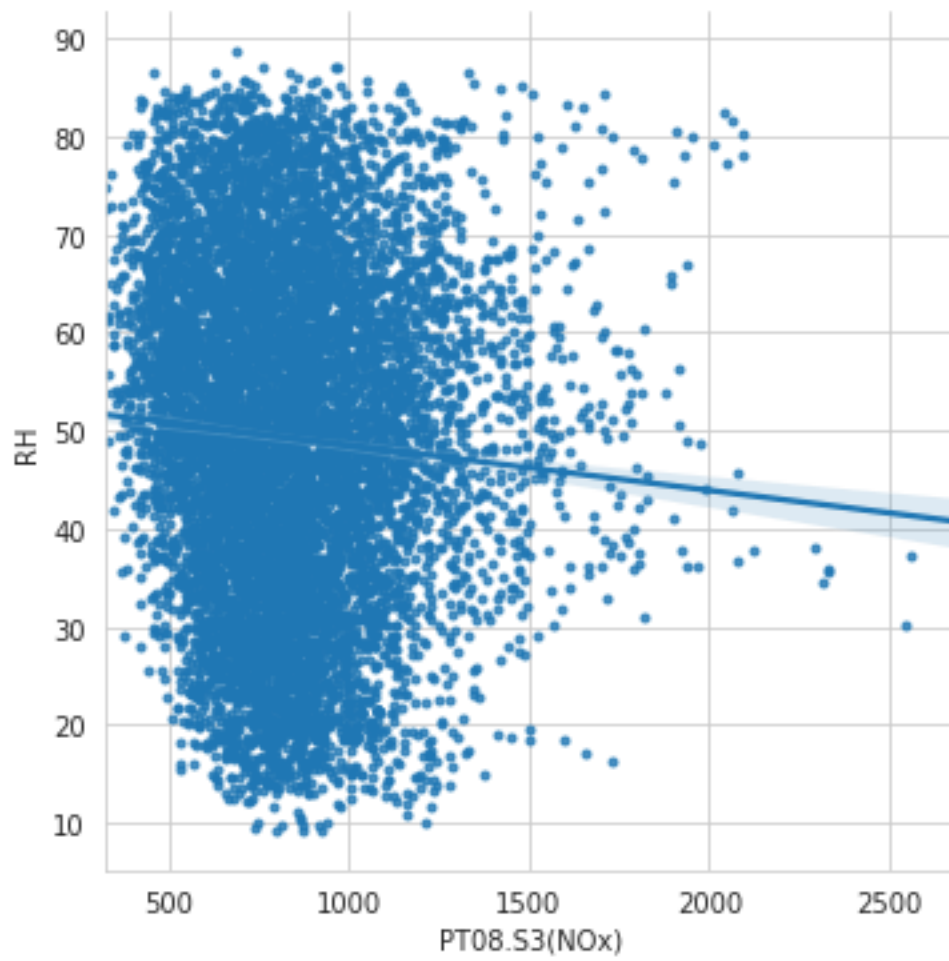


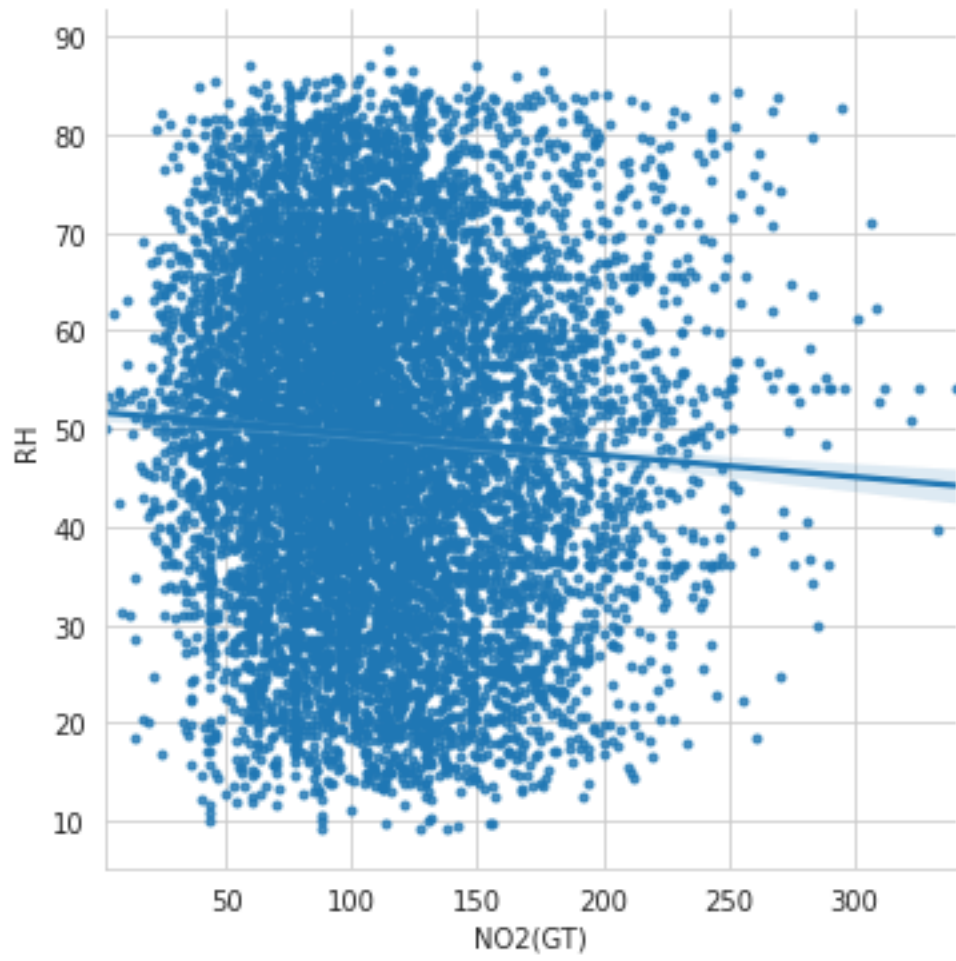








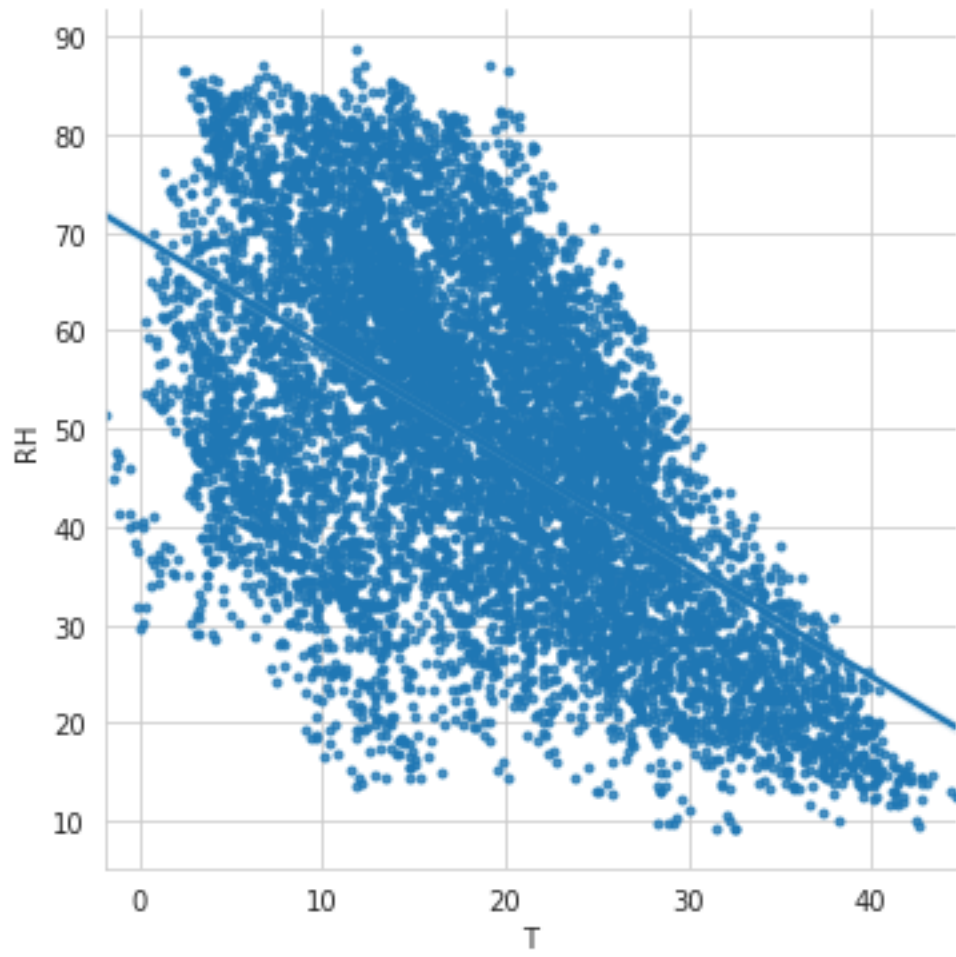


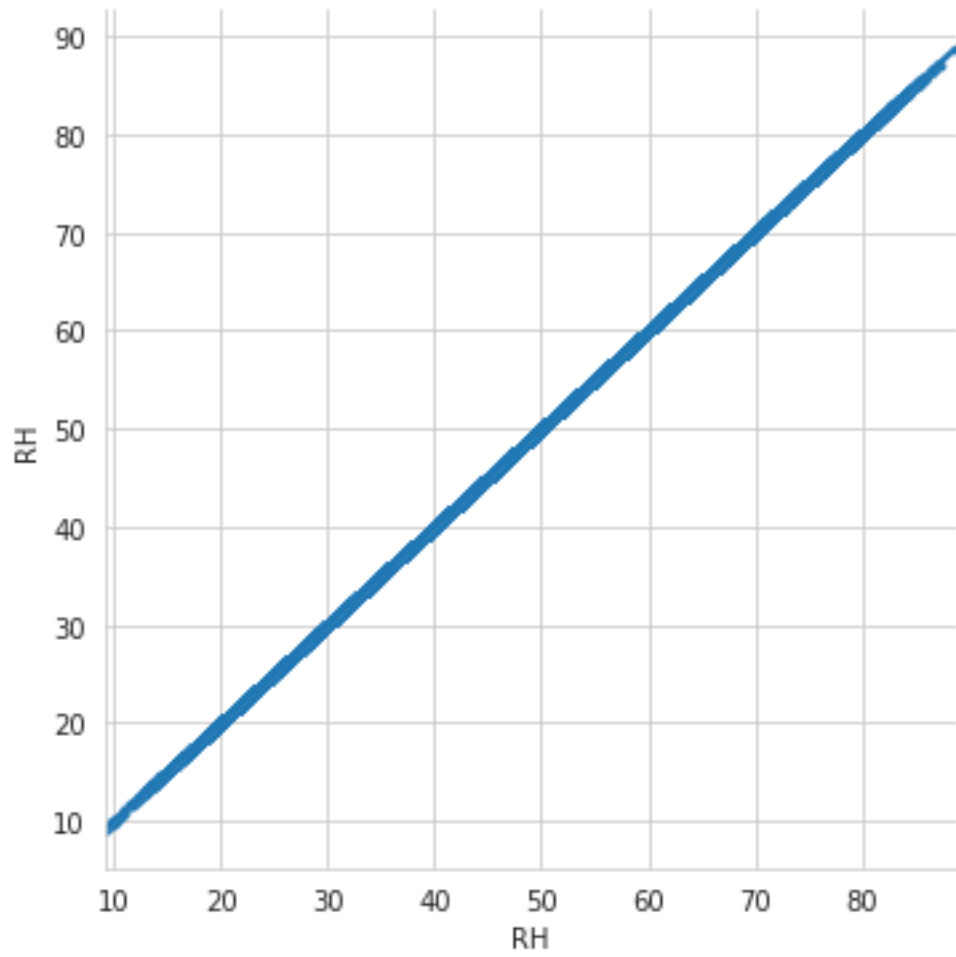


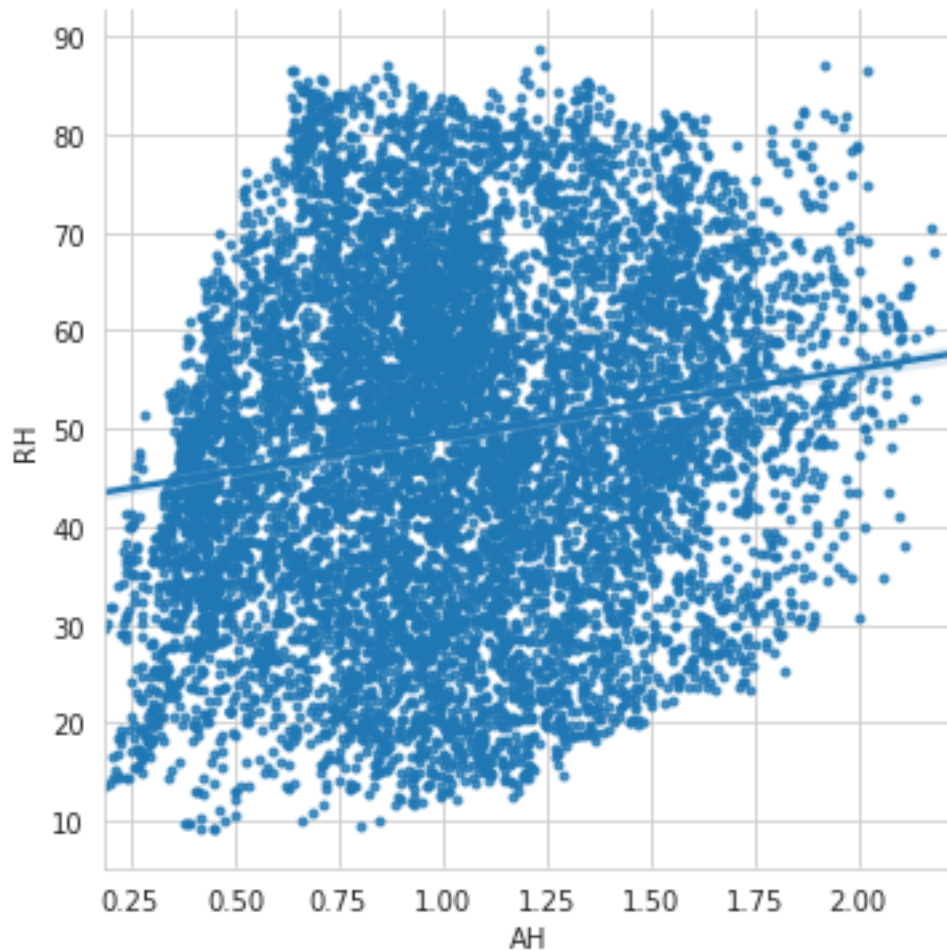












### 1.5 ML Linear, Lasso, Ridge, Elastic Net models

```
[46]: X = airq.drop(['NOx(GT)', 'T', 'Time'], axis=1)
```

```
y= airq['NOx(GT)']
```

```
[47]: X = np.log(X)
y = np.log(y)
```

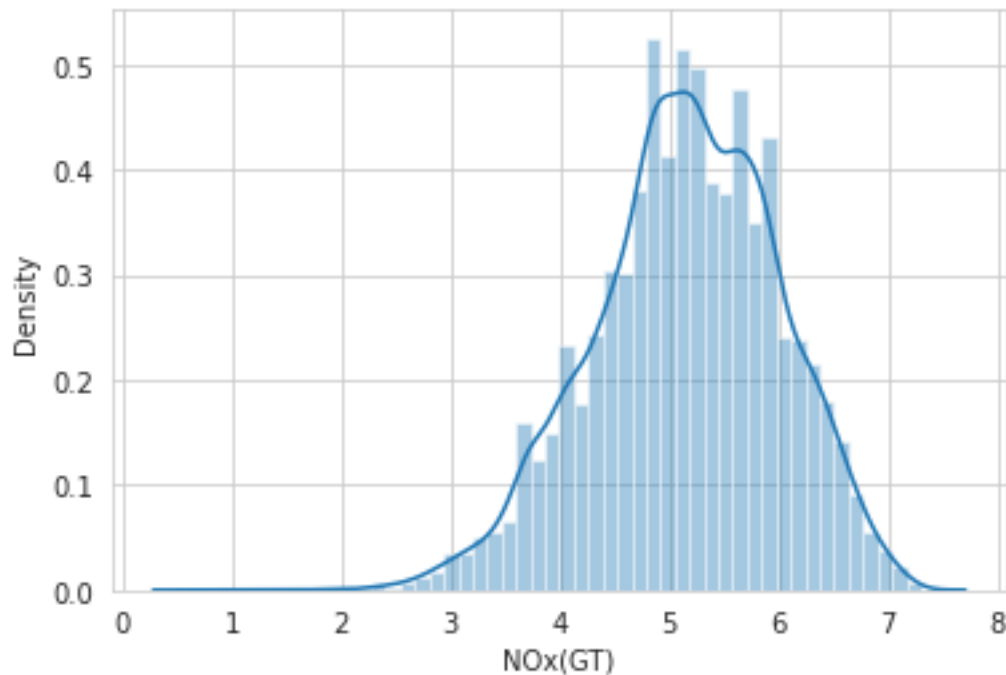
```
[48]: sns.distplot(y)
```

/home/tselest/.local/lib/python3.9/site-packages/seaborn/distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



```
[48]: <AxesSubplot:xlabel='NOx(GT)', ylabel='Density'>
```



```
[49]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                    random_state=101)
```

```
[50]: # Show the results of the split
print("Training set has {} samples.".format(X_train.shape[0]))
print("Testing set has {} samples.".format(X_test.shape[0]))
```

Training set has 6549 samples.

Testing set has 2808 samples.

```
[62]: kf = KFold(shuffle=True, random_state=72018, n_splits=3)
```

```
s = StandardScaler()
lr = LinearRegression()

X_train = s.fit_transform(X_train)
X_test = s.transform(X_test)
lr.fit(X_train, y_train)

y_pred = lr.predict(X_test)
score = r2_score(y_test, y_pred)

# with pipeline
```

```
estimator = Pipeline([("scaler", s), ("regression", lr)])
predictions_lr = cross_val_predict(estimator, X_train, y_train, cv=kf)
linear_score = r2_score(y_train, predictions_lr)


linear_score, score #almost identical
```

[62]: (0.8773345173587086, 0.8765640967999826)

```
[63]: sns.scatterplot(y_test, y_pred, alpha = 0.2)
plt.xlabel('Actual Values')
plt.ylabel('Predictions')
```

/home/tselest/.local/lib/python3.9/site-packages/seaborn/\_decorators.py:36:  
FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

[63]: Text(0, 0.5, 'Predictions')



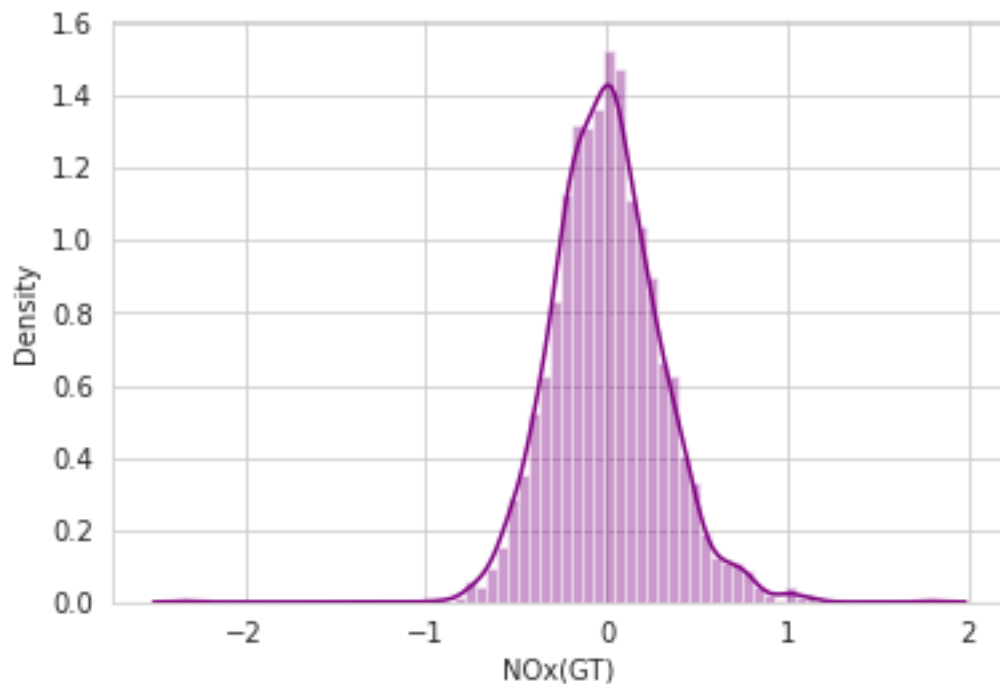
```
[72]: sns.distplot((y_test-y_pred), bins=70, color="purple")
```

/home/tselest/.local/lib/python3.9/site-packages/seaborn/distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level

function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
[72]: <AxesSubplot:xlabel='NOx(GT)', ylabel='Density'>
```



```
[64]: print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
MAE: 0.2290190551198133
MSE: 0.08797015944419892
RMSE: 0.2965976389727317
```

```
[65]: # lasso regression and K-fold cross validation
pf = PolynomialFeatures(degree=3)
kf = KFold(shuffle=True, random_state=72018, n_splits=3)
scores = []
alphas = np.geomspace(0.06, 6.0, 20)
predictions_lsr = []
for alpha in alphas:
    las = Lasso(alpha=alpha, max_iter=100000)

    estimator = Pipeline([
        ("scaler", s),
```

```

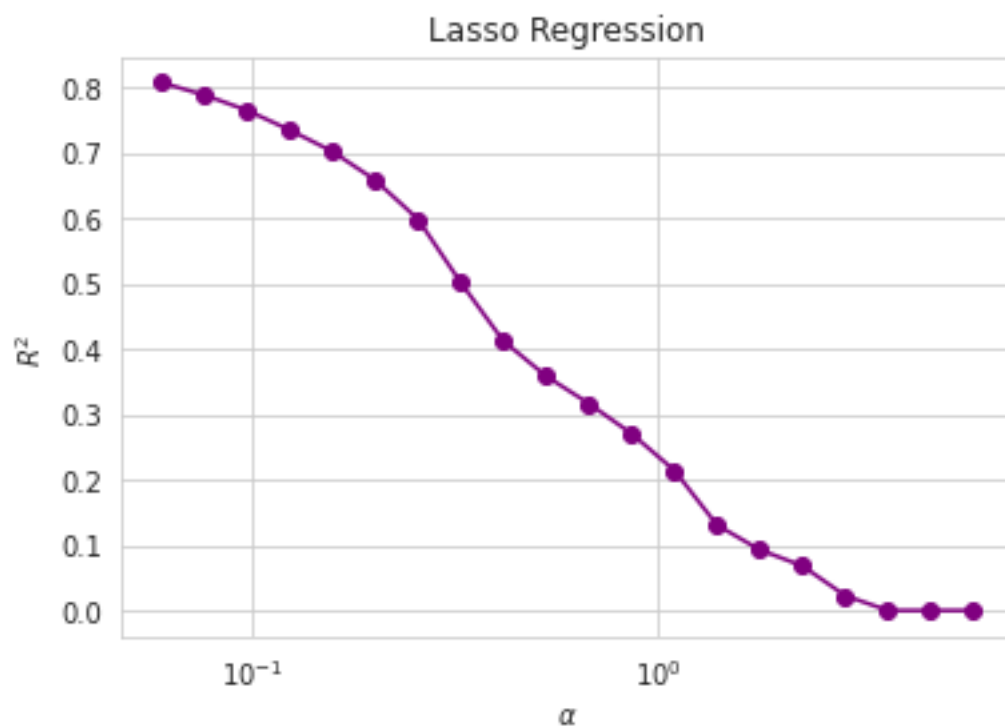
    ("make_higher_degree", pf),
    ("lasso_regression", las)])

predictions_lsr = cross_val_predict(estimator, X_train, y_train, cv = kf)

score = r2_score(y_train, predictions_lsr)

scores.append(score)
plt.semilogx(alphas, scores, '-o', color='purple')
plt.title('Lasso Regression')
plt.xlabel('$\alpha$')
plt.ylabel('$R^2$');

```



```

[66]: best_estimator = Pipeline([
        ("scaler", s),
        ("make_higher_degree", PolynomialFeatures(degree=2)),
        ("lasso_regression", Lasso(alpha=0.03))])

best_estimator.fit(X_train, y_train)
lasso_score = best_estimator.score(X_train, y_train)

```

```

[67]: # ridge regression and K-fold cross validation
pf = PolynomialFeatures(degree=2)

```

```

alphas = np.geomspace(4, 20, 20)
scores=[]
predictions_rr = []
for alpha in alphas:
    ridge = Ridge(alpha=alpha, max_iter=100000)

    estimator = Pipeline([
        ("scaler", s),
        ("polynomial_features", pf),
        ("ridge_regression", ridge)])

    predictions_rr = cross_val_predict(estimator, X_train, y_train, cv = kf)
    score = r2_score(y_train, predictions_rr)
    scores.append(score)

plt.plot(alphas, scores, '-o', color='purple')
plt.title('Ridge Regression')
plt.xlabel('$\alpha$')
plt.ylabel('$R^2$');

```



```

[68]: best_estimator = Pipeline([
        ("scaler", s),
        ("make_higher_degree", PolynomialFeatures(degree=2)),

```

```

        ("ridge_regression", Ridge(alpha=0.03))])

best_estimator.fit(X_train, y_train)
ridge_score = best_estimator.score(X_train, y_train)

```

```

[69]: # comparing accuracy scores
pd.DataFrame([[linear_score, lasso_score, ridge_score]], columns=['linear', '
↳ 'lasso', 'ridge'], index=['score'])

```

```

[69]:          linear      lasso      ridge
score  0.877335  0.847026  0.924211

```

```

[70]: def rmse(ytrue, ypredicted):
        return np.sqrt(mean_squared_error(ytrue, ypredicted))

# Fit a basic linear regression model
linearRegression = LinearRegression().fit(X_train, y_train)
linearRegression_rmse = rmse(y_test, linearRegression.predict(X_test))

# Fit a regular (non-cross validated) Ridge model
alphas = [0.005, 0.05, 0.1, 0.3, 1, 3, 5, 10, 15, 30, 80]
ridgeCV = RidgeCV(alphas=alphas, cv=4).fit(X_train, y_train)
ridgeCV_rmse = rmse(y_test, ridgeCV.predict(X_test))

# Fit a Lasso model using cross validation and determine the optimum value for
alphas2 = np.array([1e-5, 5e-5, 0.0001, 0.0005])
lassoCV = LassoCV(alphas=alphas2,
                  max_iter=5e4,
                  cv=3).fit(X_train, y_train)
lassoCV_rmse = rmse(y_test, lassoCV.predict(X_test))

# Fit elastic net with the same set of alphas as lasso
l1_ratios = np.linspace(0.1, 0.9, 9)
elasticNetCV = ElasticNetCV(alphas=alphas2,
                             l1_ratio=l1_ratios,
                             max_iter=1e4).fit(X_train, y_train)
elasticNetCV_rmse = rmse(y_test, elasticNetCV.predict(X_test))

rmse_vals = [linearRegression_rmse, ridgeCV_rmse, lassoCV_rmse, '
↳ elasticNetCV_rmse]

labels = ['Linear', 'Lasso', 'Ridge', 'ElasticNet']

# creating a pandas dataframe for comparing root-mean square errors

```

```
rmse_df = pd.DataFrame([[linearRegression_rmse, ridgeCV_rmse, lassoCV_rmse,
↳elasticNetCV_rmse]],columns=['Linear', 'Lasso', 'Ridge', 'ElasticNet'],
↳index=['rmse'])
rmse_df
```

```
[70]:
```

	Linear	Lasso	Ridge	ElasticNet
rmse	0.296598	0.296621	0.296621	0.29665

```
[71]: # plotting the results: prediction vs actual values

f = plt.figure(figsize=(6,6))
ax = plt.axes()

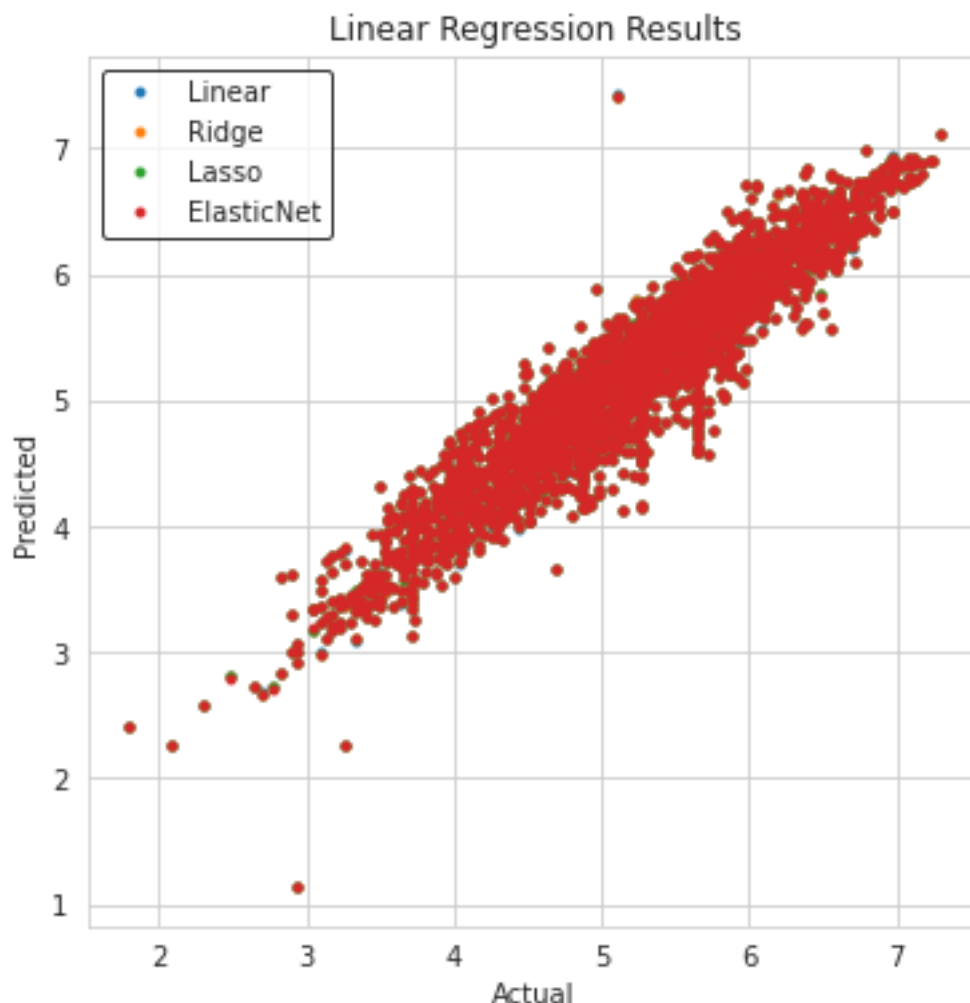
labels, models = ['Linear', 'Ridge', 'Lasso', 'ElasticNet'], [linearRegression,
↳ridgeCV, lassoCV, elasticNetCV]

for mod, label in zip(models, labels):
    ax.plot(y_test, mod.predict(X_test), marker='o', ls='', ms=3.0,
↳label=label, alpha=1)

leg = plt.legend(frameon=True)
leg.get_frame().set_edgecolor('black')
leg.get_frame().set_linewidth(1.0)

ax.set(xlabel='Actual', ylabel='Predicted', title='Linear Regression Results')
```

```
[71]: [Text(0.5, 0, 'Actual'),
Text(0, 0.5, 'Predicted'),
Text(0.5, 1.0, 'Linear Regression Results')]
```



## 1.6 Conclusion

Two of the most toxicologically significant compounds are nitric oxide (NO) and nitrogen dioxide (NO<sub>2</sub>). Nitrogen Oxides (NO<sub>x</sub>) are among the most dangerous forms of air pollution. They are produced from the reaction of nitrogen and oxygen gases in the air during combustion, especially at high temperatures. In areas of high motor vehicle traffic, such as in large cities, the amount of nitrogen oxides emitted into the atmosphere as air pollution can be significant. It is mainly due to fossil fuel combustion from both stationary sources, i.e. power generation (21%), and mobile sources, i.e. transport (44%). Other atmospheric contributions come from non-combustion processes, for example nitric acid manufacture, welding processes and the use of explosives.

In addition, these create serious health issues. These mainly impact on respiratory conditions causing inflammation of the airways at high levels. Long term exposure can decrease lung function, increase the risk of respiratory conditions and increases the response to allergens. NO<sub>x</sub> also contributes to the formation of fine particles (PM) and ground level ozone, both of which are associated with adverse health effects.



In this project, the AirQuality dataset from UCI was used, for exploratory data analysis and the prediction of NO<sub>x</sub>. linear regression models (Linear baseline, ridge, lasso and elasticnet) were created and trained, using the same training and test splits, and then compared to find the best model among them.

Based on the models findings, the simple linear model gives the smallest Root-mean-square error. However, the difference in scores and errors are not significant and almost identical. Therefore it is recommended as a final model as it best fits the data in terms of accuracy.

The above models could give even better results if we used GridSearchCV or RandomizedSearchCV to optimize the models' hyperparameters. Alternatively, different techniques like Random Forest Regression or Support Vector Machine could be used.

Here, NO<sub>x</sub> levels were predicted, using the Air Quality dataset. However, there are many others ways to measure air pollution, including PM<sub>10</sub> (particulate matter around between 2.5 and 10 microns in diameter), carbon monoxide, sulfur dioxide, nitrogen dioxide, ozone (O<sub>3</sub>), etc.